


# Introduction to CNN

Nguyen Quoc Thai



# CONTENT

**(1) – Neural Network**

**(2) – Convolutional Layer**

**(3) – Pooling Layer**

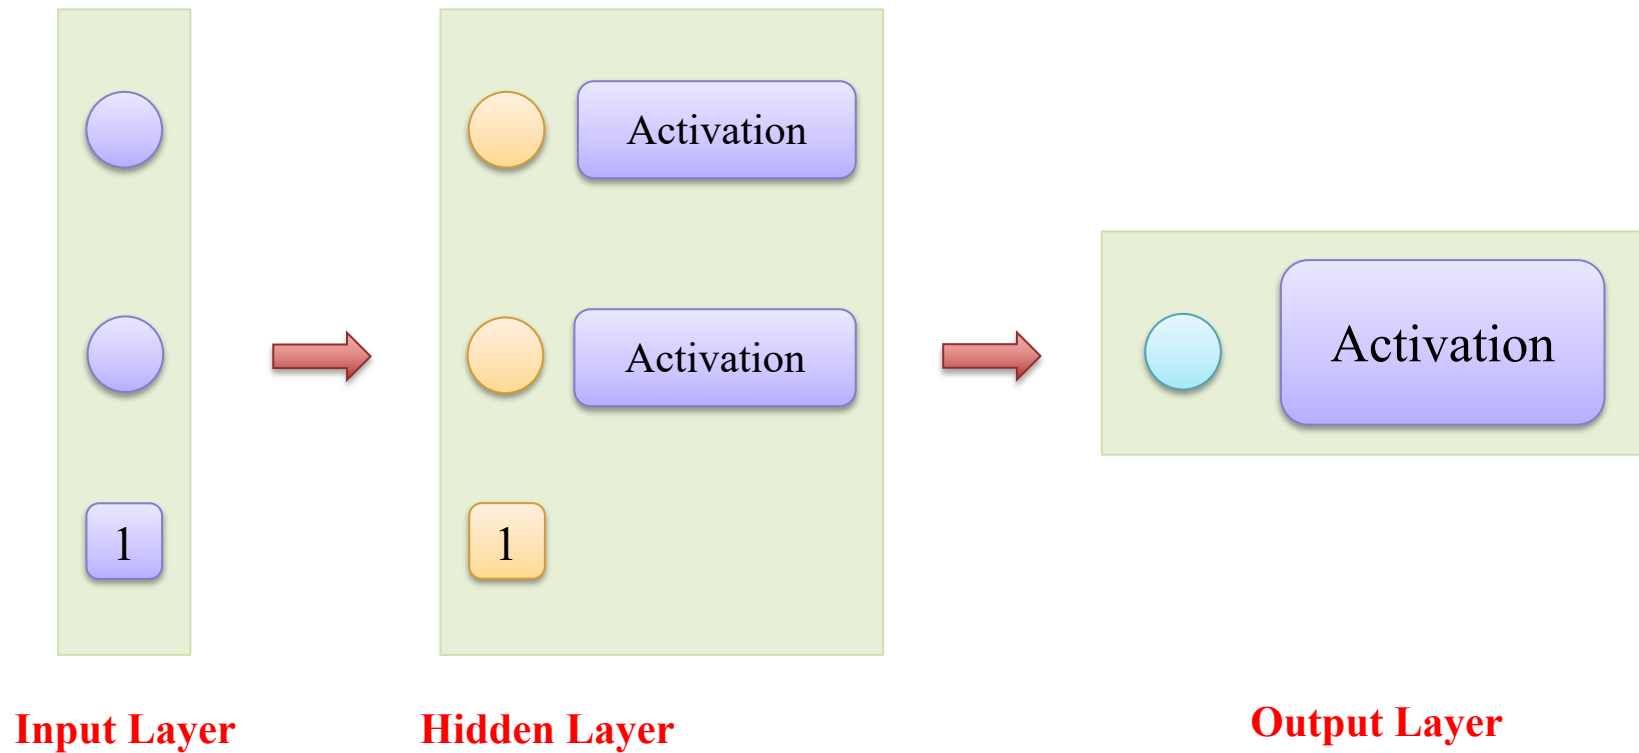
**(4) – Flatten**

**(5) – Practice**

# 1 – Neural Network



## Neural Network



Loss: CrossEntropyLoss

$$L(\theta) = - \sum_i y_i \log(\hat{y}_i)$$

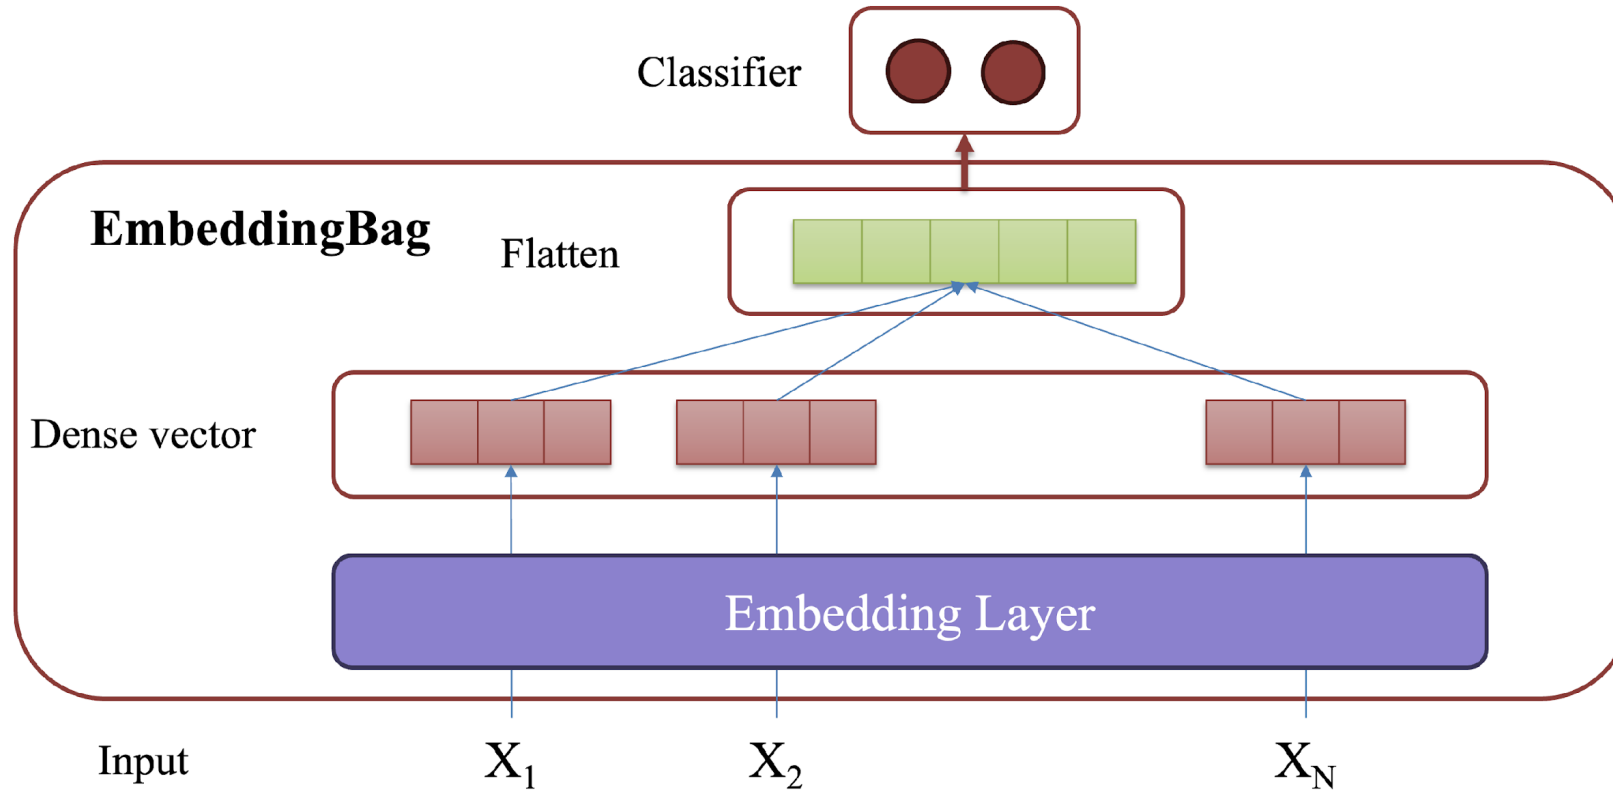
Optimizer: SGD

$$x = x - \eta * f'(x)$$

# 1 – Neural Network

## Neural Network for Text (Time Series)

- ❖ No capture the order and importance of words in a sentence

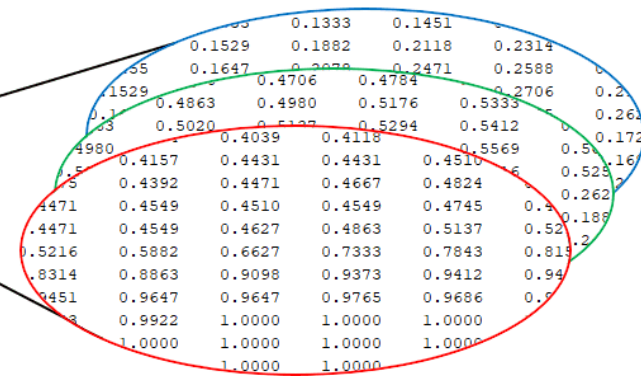




# 1 – Neural Network




## Neural Network for Image

- ❖ Each hidden node connects to all the other nodes





# 1 – Neural Network



## Neural Network

❖ **Need better network architectures...**

RNNs for Sequence

CNNs for Image


# 2 – Convolutional Layer



## Convolutional Operation

### ❖ Element-wise Multiplication Matrix

📌  $A (M \times N) \times B (M \times N) \Rightarrow C (M \times N)$



# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

=

9			

Output: 4 x 4



# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

=

9	13		

Output: 4 x 4

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

=

9	13	9	13
14	11	13	10
12	17	11	14
12	13	13	18

Output: 4 x 4

# 2 – Convolutional Layer



## Convolutional Operation

### ❖ Pytorch

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
```

```
tensor([[[[2., 2., 1., 4., 1., 0.],
          [0., 4., 0., 3., 3., 4.],
          [0., 4., 1., 2., 0., 0.],
          [2., 1., 4., 1., 3., 1.],
          [4., 3., 1., 4., 2., 4.],
          [2., 0., 0., 4., 3., 4.]]]]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    bias=False
)
```

```
conv_layer.weight
```

```
Parameter containing:
tensor([[[[ 0.0520,  0.2693,  0.0364],
          [-0.1051,  0.0896, -0.0904],
          [ 0.1403,  0.2976,  0.1927]]]], requires_grad=True)
```

```
init_kernel_weight = torch.randint(
    high=2,
    size=(conv_layer.weight.data.shape),
    dtype=torch.float32
)
init_kernel_weight
```

```
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]]]])
```

```
# init weight
conv_layer.weight.data = init_kernel_weight
```

```
conv_layer.weight
```

```
Parameter containing:
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]]], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[[ 9., 13.,  9., 13.],
          [14., 11., 13., 10.],
          [12., 17., 11., 14.],
          [12., 13., 13., 18.]]]], grad_fn=<SqueezeBackward1>)
```

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

1
---

Bias

=

10			

Output: 4 x 4

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	1	0
1	1	0
1	0	1

Kernel: 3 x 3

1

Bias

=

10	14	10	14
15	12	14	11
13	18	12	15
13	14	14	19

Output: 4 x 4

# 2 – Convolutional Layer



## Convolutional Operation

### ❖ Pytorch

input

```
tensor([[[[2., 2., 1., 4., 1., 0.],
          [0., 4., 0., 3., 3., 4.],
          [0., 4., 1., 2., 0., 0.],
          [2., 1., 4., 1., 3., 1.],
          [4., 3., 1., 4., 2., 4.],
          [2., 0., 0., 4., 3., 4.]]]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
)
```

init\_kernel\_weight

```
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]]])
```

```
# init weight
conv_layer.weight.data = init_kernel_weight
```

conv\_layer.weight

```
Parameter containing:
tensor([[[[1., 1., 0.],
          [1., 1., 0.],
          [1., 0., 1.]]]], requires_grad=True)
```

conv\_layer.bias

```
Parameter containing:
tensor([-0.1148], requires_grad=True)
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
```

conv\_layer.bias

```
Parameter containing:
tensor([1.], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[[10., 14., 10., 14.],
          [15., 12., 14., 11.],
          [13., 18., 12., 15.],
          [13., 14., 14., 19.]]]], grad_fn=<SqueezeBackward1>)
```

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

1	0	1
0	1	1

Kernel: 2 x 3

\*

=

1
---

Bias

8			

Output: 5 x 4

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input: 6 x 6

\*

1	0	1
0	1	1

Kernel: 2 x 3

=

1
---

Bias

8	10	9	12
6	11	6	8
7	12	6	7
11	8	14	9
6	12	11	16

Output: 5 x 4



# 2 – Convolutional Layer



## Convolutional Operation

### ❖ Pytorch

input

```
tensor([[[2., 2., 1., 4., 1., 0.],
         [0., 4., 0., 3., 3., 4.],
         [0., 4., 1., 2., 0., 0.],
         [2., 1., 4., 1., 3., 1.],
         [4., 3., 1., 4., 2., 4.],
         [2., 0., 0., 4., 3., 4.]])])
```

```
init_kernel_weight = torch.randint(
    high=2,
    size=(conv_layer.weight.data.shape),
    dtype=torch.float32
)
```

init\_kernel\_weight

```
tensor([[[[1., 0., 1.],
         [0., 1., 1.]])])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=(2, 3), # create a kernel: 2 x 3
)
```

```
# init weight & bias
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

Parameter containing:  
tensor([[[[1., 0., 1.],  
 [0., 1., 1.]])], requires\_grad=True)

conv\_layer.bias

Parameter containing:  
tensor([0.3672], requires\_grad=True)

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

Parameter containing:  
tensor([1.], requires\_grad=True)

output = conv\_layer(input)  
output

```
tensor([[[ 8., 10., 9., 12.],
         [ 6., 11., 6., 8.],
         [ 7., 12., 6., 7.],
         [11., 8., 14., 9.],
         [ 6., 12., 11., 16.]])], grad_fn=<SqueezeBackward1>)
```

# 2 – Convolutional Layer



## Convolutional Operation

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1
4	3	1	4	2	4
2	0	0	4	3	4

Input:  $M \times N$

1	0	1
0	1	1

Kernel:  $K \times O$

\*

=


1
---

Bias

8	10	9	12
6	11	6	8
7	12	6	7
11	8	14	9
6	12	11	16

Output:

$M - (K - 1) \times N - (O - 1)$



# 2 – Convolutional Layer



## Padding

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input: 4 x 4

Padding: 1 x 1

0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0

Shape: 6 x 6

\*

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1
---

Bias

=

7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6

Output: 4 x 4

# 2 – Convolutional Layer



## Padding

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[[2., 3., 1., 4.],
          [1., 1., 3., 2.],
          [0., 4., 3., 0.],
          [3., 2., 2., 0.]]]])
```

```
init_kernel_weight = torch.randint(
    high=2,
    size=(conv_layer.weight.data.shape),
    dtype=torch.float32
)
init_kernel_weight
```

```
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding='same'
)
```

{“valid”, “same”}

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```


```
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires_grad=True)
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
tensor([1.], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[[ 7.,  8., 12.,  8.],
          [ 8., 16., 18., 11.],
          [10., 15., 16.,  9.],
          [10., 15., 12.,  6.] ]]]], grad_fn=<SqueezeBackward1>)
```



# 2 – Convolutional Layer



## Padding

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input: 4 x 4

Padding: 2 x 1

0	0	0	0	0	0
0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Shape: 8 x 6

\*

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1
---

Bias

=

3	4	2	5
7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6
6	8	5	3

Output: 6 x 4

# 2 – Convolutional Layer



## Padding

input

```
tensor([[[2., 3., 1., 4.],
         [1., 1., 3., 2.],
         [0., 4., 3., 0.],
         [3., 2., 2., 0.]]])
```

```
# define convolutional layer
```

```
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding=(2, 1)
)
```

An int / a tuple of ints

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

Parameter containing:

```
tensor([[[[1., 1., 1.],
         [1., 1., 1.],
         [0., 1., 0.]]]], requires_grad=True)
```

```
# init bias
```

```
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

Parameter containing:

```
tensor([1.], requires_grad=True)
```

```
output = conv_layer(input)
```

output

```
tensor([[[ 3.,  4.,  2.,  5.],
         [ 7.,  8., 12.,  8.],
         [ 8., 16., 18., 11.],
         [10., 15., 16.,  9.],
         [10., 15., 12.,  6.],
         [ 6.,  8.,  5.,  3.] ]], grad_fn=<SqueezeBackward1>)
```

# 2 – Convolutional Layer

!

## Padding

2	3	1	4
1	1	3	2
0	4	3	0
3	2	2	0

Input:  $M \times N$ Padding:  $P \times Q$ 

0	0	0	0	0	0
0	0	0	0	0	0
0	2	3	1	4	0
0	1	1	3	2	0
0	0	4	3	0	0
0	3	2	2	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Shape:  $(M+2P) \times (N+2Q)$ 

\*

1	1	1
1	1	1
0	1	0

Kernel:  $K \times O$ 

1

Bias

=

3	4	2	5
7	8	12	8
8	16	18	11
10	15	16	9
10	15	12	6
6	8	5	3

Output:

 $(M+2P-K+1) \times (N+2Q-O+1)$

# 2 – Convolutional Layer



## Stride

Stride: 1 (1x1)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

\*

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1

Bias

=

10	10	13	15
10	12	14	15
11	12	16	17
12	16	14	16

Output: 4 x 4



# 2 – Convolutional Layer



## Stride

Stride: 2 (2x2)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

\*

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1
---

Bias

=

10	

Output: 2 x 2

# 2 – Convolutional Layer

!

## Stride

		Skip		Skip	
1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

Stride: 2 (2x2)

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1

Bias

\*

=

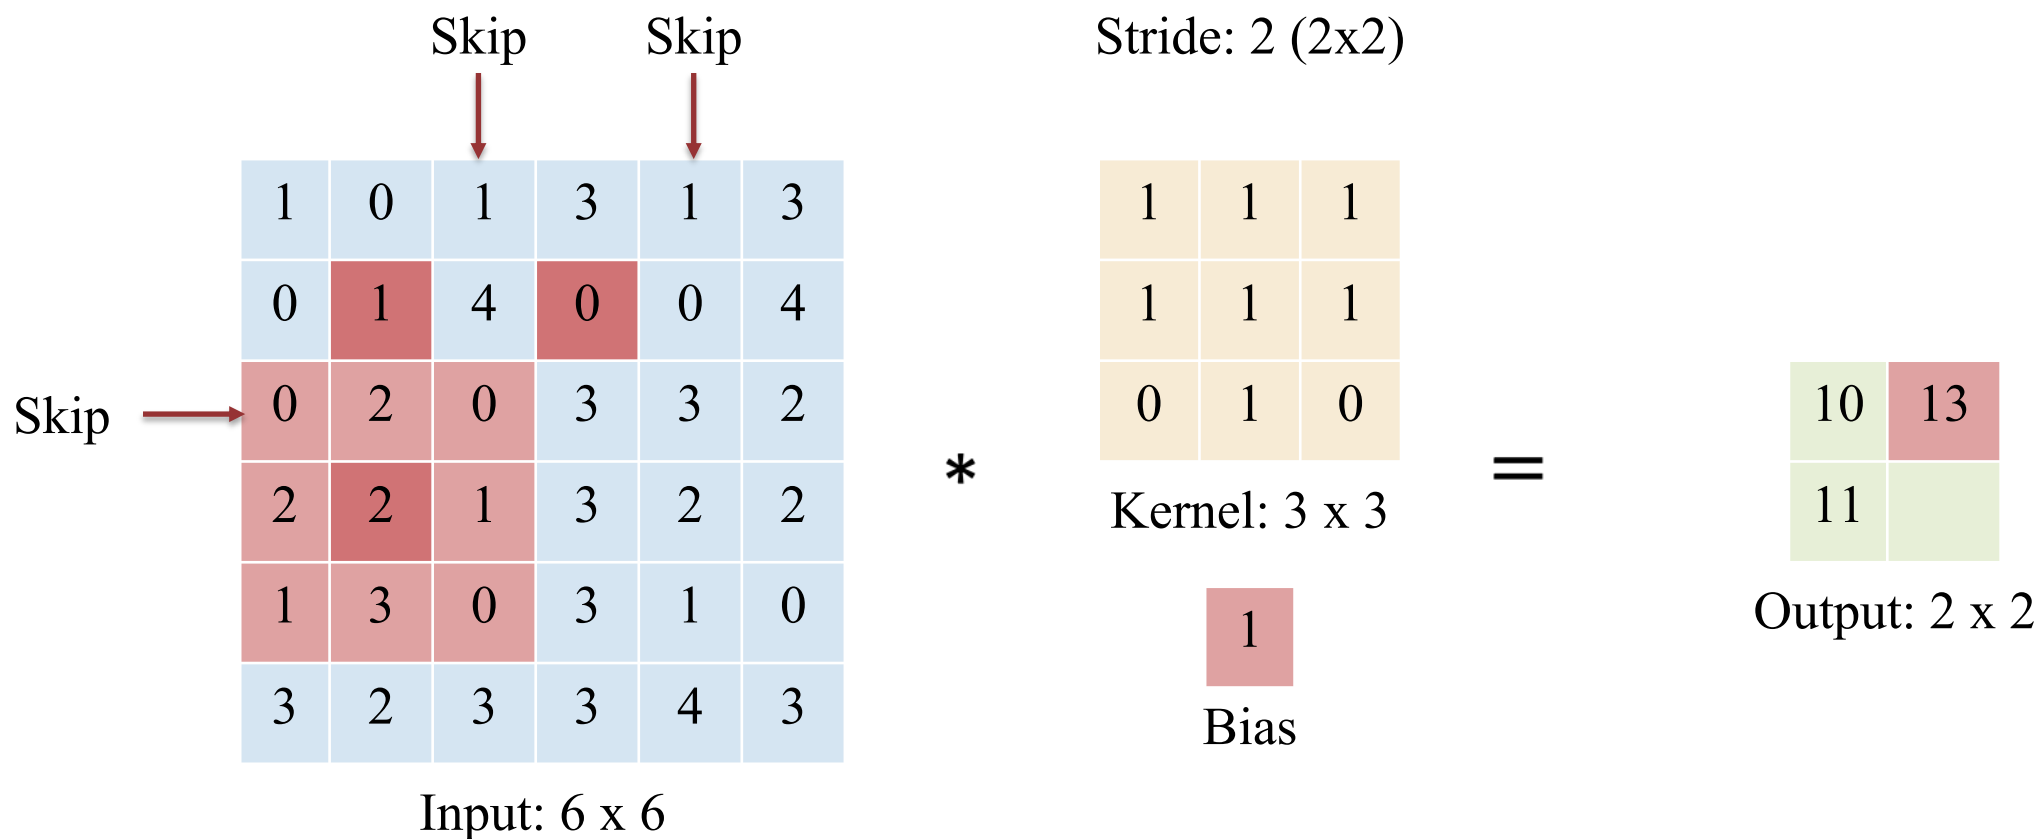
10	13

Output: 2 x 2

# 2 – Convolutional Layer



## Stride



# 2 – Convolutional Layer



## Stride

Stride: 2 (2x2)

1	0	1	3	1	3
0	1	4	0	0	4
0	2	0	3	3	2
2	2	1	3	2	2
1	3	0	3	1	0
3	2	3	3	4	3

Input: 6 x 6

\*

1	1	1
1	1	1
0	1	0

Kernel: 3 x 3

1
---

Bias

=

10	13
11	16

Output: 2 x 2

# 2 – Convolutional Layer



## Stride

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
```

```
tensor([[[[1., 0., 1., 3., 1., 3.],
          [0., 1., 4., 0., 0., 4.],
          [0., 2., 0., 3., 3., 2.],
          [2., 2., 1., 3., 2., 2.],
          [1., 3., 0., 3., 1., 0.],
          [3., 2., 3., 3., 4., 3.]]]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    stride=2
)
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```


```
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires_grad=True)
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
tensor([1.], requires_grad=True)
```

```
output = conv_layer(input)
output
```

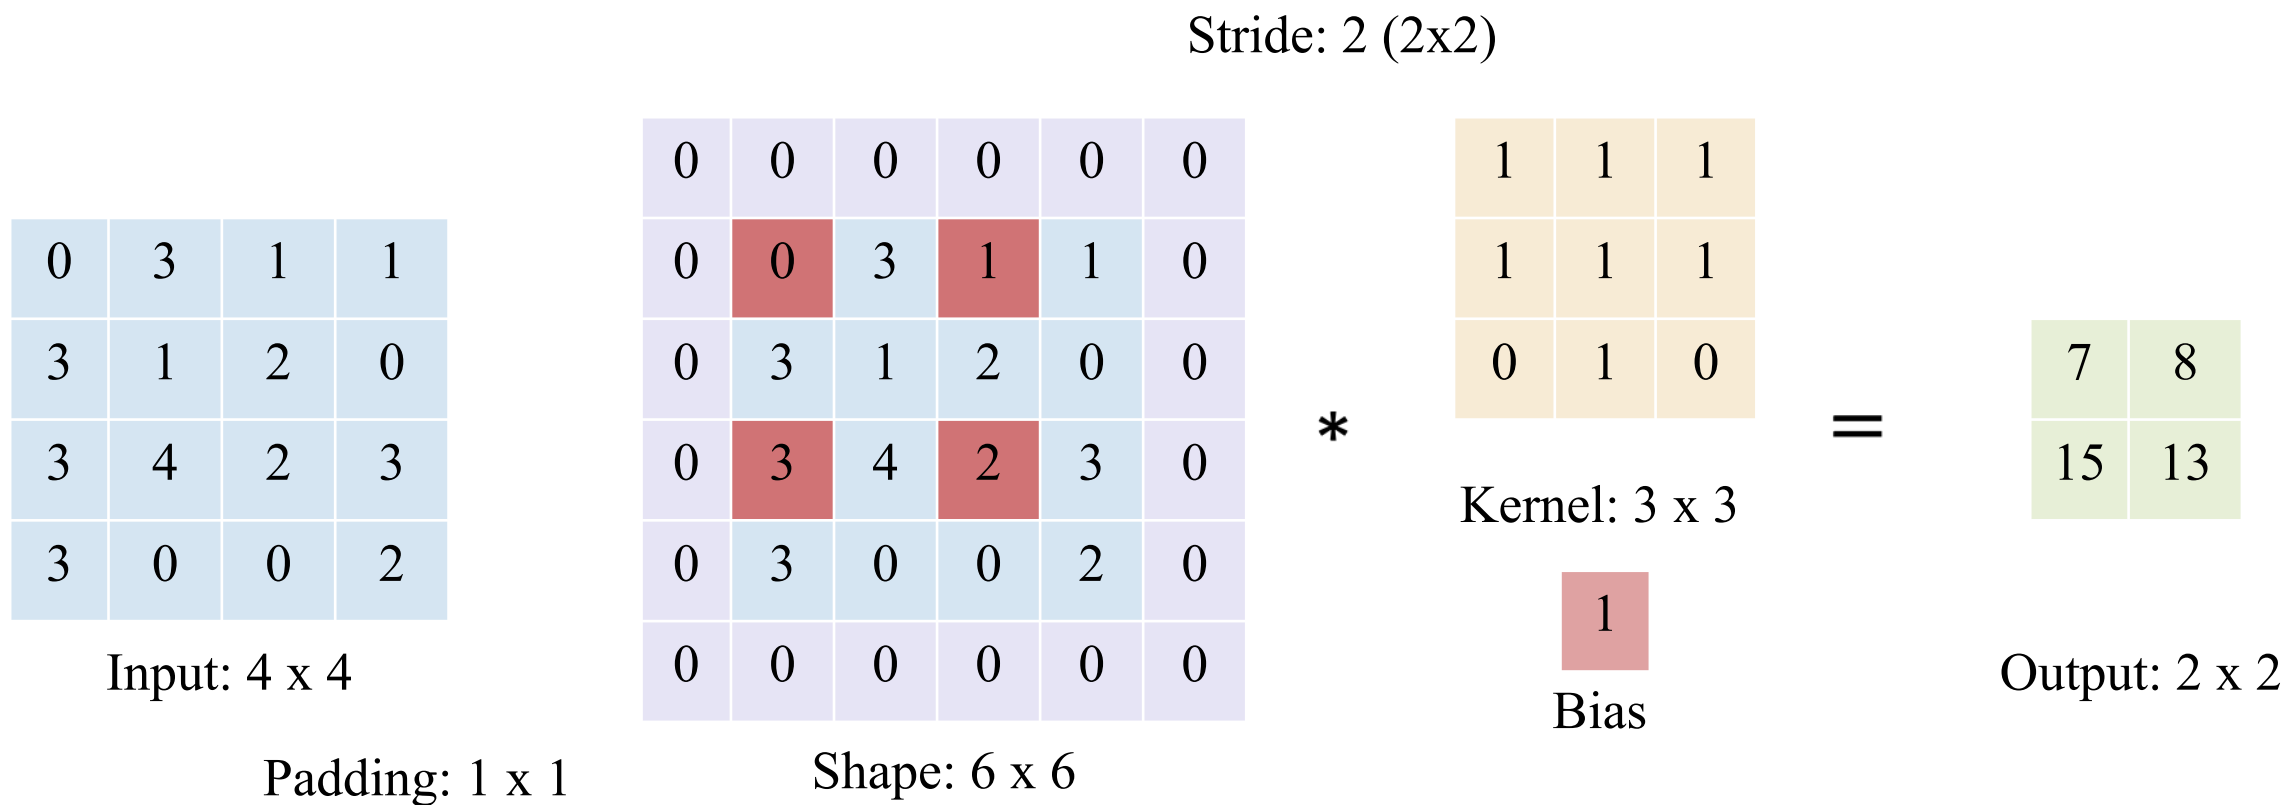
```
tensor([[[[10., 13.],
          [11., 16.]]]], grad_fn=<SqueezeBackward1>)
```



# 2 – Convolutional Layer



## Stride



# 2 – Convolutional Layer



## Stride

```
input = torch.randint(5, (1, 4, 4), dtype=torch.float32)
input
```

```
tensor([[[[0., 3., 1., 1.],
          [3., 1., 2., 0.],
          [3., 4., 2., 3.],
          [3., 0., 0., 2.]]]])
```

```
# define convolutional layer
conv_layer = nn.Conv2d(
    in_channels=1,
    out_channels=1,
    kernel_size=3, # create a kernel: 3 x 3
    padding=1,
    stride=(2, 2)
)
```

```
conv_layer.weight.data = init_kernel_weight
conv_layer.weight
```

```
Parameter containing:
tensor([[[[1., 1., 1.],
          [1., 1., 1.],
          [0., 1., 0.]]]], requires_grad=True)
```

```
# init bias
conv_layer.bias = nn.Parameter(
    torch.tensor([1], dtype=torch.float32)
)
conv_layer.bias
```

```
Parameter containing:
tensor([1.], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[ 7.,  8.],
          [15., 13.] ]], grad_fn=<SqueezeBackward1>)
```

# 2 – Convolutional Layer



## Stride

Stride: (S, T)

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

Input: M x N

Padding: (P, Q)

0	0	0	0	0	0
0	0	3	1	1	0
0	3	1	2	0	0
0	3	4	2	3	0
0	3	0	0	2	0
0	0	0	0	0	0

Shape: (M+2P) x (N+2Q)

\*

1	1	1
1	1	1
0	1	0

Kernel: K x O

1


Bias

=

7	8
15	13

$$\left\lfloor \frac{M + 2P - K}{S} + 1 \right\rfloor \times \left\lfloor \frac{N + 2Q - K}{T} + 1 \right\rfloor$$





# 3 – Pooling Layer



## Max Pooling

Kernel Size: 2  
Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6


3	2
0	3

Max values


=

3		

Output: 3 x 3



# 3 – Pooling Layer



## Max Pooling

Kernel Size: 2  
Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6


1	0
3	1

Max values

=

3	3	

Output: 3 x 3



# 3 – Pooling Layer



## Max Pooling

Kernel Size: 2  
Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6


1	0
3	1

Max values

=

3	3	3
4	4	4
4	4	4

Output: 3 x 3



# 3 – Pooling Layer



## Max Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2  
Stride: 2

3	3	3
4	4	4
4	4	4

Output: 3 x 3

```
input = torch.randint(5, (1, 6, 6), dtype=torch.float32)
input
```

```
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]]])
```

```
max_pool_layer = nn.MaxPool2d(kernel_size=2)
```

```
output = max_pool_layer(input)
output
```

Default: Stride = 2

```
tensor([[[3., 3., 3.],
         [4., 4., 4.],
         [4., 4., 4.]])])
```

# 3 – Pooling Layer



## Max Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2  
Stride: (1, 2)

3	3	4
3	4	1
4	4	4
4	3	4
4	4	4

Output: 5 x 3

input

```
tensor([[[[3., 2., 1., 0., 0., 3.],
           [0., 3., 3., 1., 1., 0.],
           [3., 1., 4., 1., 1., 0.],
           [2., 4., 1., 1., 0., 4.],
           [1., 0., 3., 0., 3., 0.],
           [3., 4., 4., 3., 3., 4.]]]])
```

```
max_pool_layer = nn.MaxPool2d(
    kernel_size=2,
    stride=(1, 2)
)
```

```
output = max_pool_layer(input)
output
```

```
tensor([[[[3., 3., 3.],
           [3., 4., 1.],
           [4., 4., 4.],
           [4., 3., 4.],
           [4., 4., 4.]]]])
```

# 3 – Pooling Layer



## Max Pooling

### MaxPool1d

Kernel Size: 3

Stride: 3

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

3	3
3	1
4	1
4	4
3	3
4	4

Output: 6 x 2

input

```
tensor([[[3., 2., 1., 0., 0., 3.],
         [0., 3., 3., 1., 1., 0.],
         [3., 1., 4., 1., 1., 0.],
         [2., 4., 1., 1., 0., 4.],
         [1., 0., 3., 0., 3., 0.],
         [3., 4., 4., 3., 3., 4.]]])
```

```
max_pool_layer = nn.MaxPool1d(
    kernel_size=3,
    stride=3
)
```

max\_pool\_layer(input)

```
tensor([[[3., 3.],
         [3., 1.],
         [4., 1.],
         [4., 4.],
         [3., 3.],
         [4., 4.]])])
```

# 3 – Pooling Layer



## Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

3	2
0	3
3	1

Average values

=

2.0		

Output: 2 x 3

# 3 – Pooling Layer



## Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

1	0
3	1
4	1

Average values

=

2	1.7	

Output: 2 x 3



# 3 – Pooling Layer



## Average Pooling

Kernel Size: (3, 2)

Stride: 2

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

3	1
2	4
1	0

Average values

=

2	1.7	0.8
1.8		

Output: 2 x 3

# 3 – Pooling Layer



## Average Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: (3, 2)  
Stride: 2

2	1.7	0.8
1.8	1.6	1.3

Output: 2 x 3


input

```
tensor([[[3., 2., 1., 0., 0., 3.],
          [0., 3., 3., 1., 1., 0.],
          [3., 1., 4., 1., 1., 0.],
          [2., 4., 1., 1., 0., 4.],
          [1., 0., 3., 0., 3., 0.],
          [3., 4., 4., 3., 3., 4.]])])
```

```
avg_pool_layer = nn.AvgPool2d(
    kernel_size=(3, 2),
    stride=(2, 2)
)
```

```
output = avg_pool_layer(input)
output
```

```
tensor([[[2.0000, 1.6667, 0.8333],
          [1.8333, 1.6667, 1.3333]])])
```



# 3 – Pooling Layer



## Average Pooling

### AvgPool1d

Kernel Size: 3

Stride: 3

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

2.0	1.0
2.0	0.7
2.7	0.7
2.3	1.7
1.3	1.0
3.7	3.3

Output: 6 x 2

input

```
tensor([[[3., 2., 1., 0., 0., 3.],  
         [0., 3., 3., 1., 1., 0.],  
         [3., 1., 4., 1., 1., 0.],  
         [2., 4., 1., 1., 0., 4.],  
         [1., 0., 3., 0., 3., 0.],  
         [3., 4., 4., 3., 3., 4.]]])
```

```
avg_pool_layer = nn.AvgPool1d(  
    kernel_size=3,  
    stride=3  
)
```

```
output = avg_pool_layer(input)  
output
```

```
tensor([[[2.0000, 1.0000],  
         [2.0000, 0.6667],  
         [2.6667, 0.6667],  
         [2.3333, 1.6667],  
         [1.3333, 1.0000],  
         [3.6667, 3.3333]])])
```

# 4 - Flatten



Flattens a contiguous range of dims into a tensor

2	4
3	1
3	4

Input: 3 x 2

2	4	3	1	3	4
---	---	---	---	---	---

Output: 1 x 6

```
input = torch.randint(5, (1, 3, 2), dtype=torch.float32)
input
```

```
tensor([[[2., 4.],
         [3., 1.],
         [3., 4.]]])
```

```
flatten_layer = nn.Flatten()
```

```
output = flatten_layer(input)
output
```

```
tensor([2., 4., 3., 1., 3., 4.])
```

# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=


# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4			

# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4	7		

# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4	7	5	



# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4	7	5	8

# 5 - Practice

1

## Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6


\*

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

=

4	7	5	8
4	8	4	8



1

Exercise – Convolutional Layer

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

1	1	0
1	0	0
0	0	0

\*      Kernel: 3 x 3      =

2  
Bias


# 5 - Practice

1

## Exercise – Convolutional Layer

Stride: 1 (1x1)

2	2	1	4	1	0
0	4	0	3	3	4
0	4	1	2	0	0
2	1	4	1	3	1

Input: 4 x 6

1	1	0
1	0	0
0	0	0

Kernel: 3 x 3

\*

2
---

Bias

=

6	9	7	10
6	10	6	10

## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 4 x 6

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 6 x 8

Stride: 1 (1x1)

\*


1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

=


2

Bias



2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 1 (1x1)

\*

1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

2

Bias

=

11	13	12

## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 1 (1x1)

\*

1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

=

11	13	12
15	18	13

2

Bias



2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 1 (1x1)

\*

1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

2

Bias

=

11	13	12
15	18	13
14	14	11



## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 1 (1x1)

\*

1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

=

11	13	12
15	18	13
14	14	11
9	17	8

2

Bias

## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 1 (1x1)

\*

1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

=

11	13	12
15	18	13
14	14	11
9	17	8
7	15	6

2

Bias

## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)

\*


1	0	1
1	1	1
0	1	0

Kernel: 3 x 3

=


2

Bias



2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5


Stride: 2 (2x2)

*	<table><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	1	0	1	1	1	1	0	1	0	=	<table><tr><td>11</td><td>12</td></tr><tr><td></td><td></td></tr><tr><td></td><td></td></tr></table>	11	12				
1	0	1																
1	1	1																
0	1	0																
11	12																	

Kernel: 3 x 3

2

Bias



2

Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)

*	<table><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr></table>	1	0	1	1	1	1	0	1	0	=	<table><tr><td>11</td><td>12</td></tr><tr><td>14</td><td>11</td></tr><tr><td></td><td></td></tr></table>	11	12	14	11		
1	0	1																
1	1	1																
0	1	0																
11	12																	
14	11																	

Kernel: 3 x 3

2

Bias

## 5 - Practice

2

## Exercise – Padding

2	4	2
3	3	4
3	2	0
4	0	4
1	4	0

Input: 5 x 3

Padding: 1 x 1

0	0	0	0	0
0	2	4	2	0
0	3	3	4	0
0	3	2	0	0
0	4	0	4	0
0	1	4	0	0
0	0	0	0	0

Input: 7 x 5

Stride: 2 (2x2)

	1	0	1	
*	1	1	1	=
	0	1	0	

11	12
14	11
7	6

Kernel: 3 x 3

2

Bias

3

## Exercise – Convolutional Layer + Pooling

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

Input: 5 x 3

Stride: 1 (1x1)

\*

1	1
1	0
0	0

Kernel: 3 x 2

=


1

Bias

## 3

## Exercise – Convolutional Layer + Pooling

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

Input: 5 x 3

Stride: 1 (1x1)

\*

1	1
1	0
0	0

Kernel: 3 x 2

=

8	10
8	8
6	4

Max Pooling  
Kernel Size: (1x2)


1

Bias



## 3

## Exercise – Convolutional Layer + Pooling

2	4	2
1	3	2
3	2	1
0	0	1
0	0	1

Input: 5 x 3

Stride: 1 (1x1)

\*

1	1
1	0
0	0

Kernel: 3 x 2

=


8	10
8	8
6	4

Max Pooling  
Kernel Size: (1x2)

10
8
6

1

Bias



4

## Exercise – Pooling For Grayscale Image


0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

Input: 7 x 7

MaxPooling  
2x2

=


Output: 3 x 3



4

## Exercise – Pooling For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0


Input: 7 x 7

MaxPooling  
2x2

=

0	43	43
191	250	251
241	198	250

Output: 3 x 3



4

## Exercise – Convolutional For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

Input: 7 x 7


\*

1	0	-1
1	0	-1
1	0	-1

Kernel: 3 x 3

=


Output: 5 x 5



4

## Exercise – Convolutional For Grayscale Image

0	0	0	0	0	0	0
0	0	0	43	43	0	0
0	30	250	230	125	251	0
0	191	38	0	0	81	0
0	241	0	35	119	250	0
0	49	193	198	83	0	0
0	0	0	0	0	0	0

Input: 7 x 7

\*


1	0	-1
1	0	-1
1	0	-1

Kernel: 3 x 3

=

-250	-243	82	22	168
-288	34	206	-59	168
212	657	294	185	244
-155	248	29	64	202
-193	127	229	486	202

Output: 5 x 5



4

Exercise – Convolutional For Grayscale Image

-250	-243	82	22	168
-288	34	206	-59	168
212	657	294	185	244
-155	248	29	64	202
-193	127	229	486	202

Input: 5 x 5

MaxPooling  
Kernel: 2

34	206
657	297



AI VIET NAM

@aivietnam.edu.vn

# Thanks!

## Any questions?