

Convolutional Neural Network

Motivation and Introduction to CNN

Quang-Vinh Dinh
Ph.D. in Computer Science

Outline

- **MLP Limitation**
- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Some Examples**
- **Application to Cifar10**

Fashion-MNIST dataset

Grayscale images

Resolution=28x28


Training set: 60000 samples

Testing set: 10000 samples

T-shirt



Trouser



Pullover




Dress




Coat




Sandal



Shirt



Sneaker



Bag




Ankle
Boot



MLP for Fashion-MNIST

Case 1

❖ ReLU, He and Adam



MLP for Fashion-MNIST

Case 1

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(784, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)
```

```
# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                               nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)
```

```
# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                        lr=0.001)
```

```
# Load CFashionMNIST dataset
transform = Compose([transforms.ToTensor(),
                     transforms.Normalize((0.5,),
                                         (0.5,))])

trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)

trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)


testset = FashionMNIST(root='data',
                      train=False,
                      download=True,
                      transform=transform)

testloader = DataLoader(testset,
                      batch_size=1024,
                      num_workers=10,
                      shuffle=False)
```


MLP for Fashion-MNIST

Case 1

❖ ReLU, He and Adam



Adam with learning rate of $1e-4$



Test Accuracy: ~89%

Perform reasonably

Cifar-10 dataset

Color images

Resolution=32x32

Training set: 50000 samples

Testing set: 10000 samples

airplane



automobile



bird



cat



deer



dog



frog



horse



ship




truck



MLP for Cifar-10

Case 2

❖ ReLU, He and Adam



MLP for Cifar-10

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(32*32*3, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                               nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                        lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                      Normalize((0.5,0.5, 0.5),
                                (0.5,0.5, 0.5))])



trainset = CIFAR10(root='data',
                   train=True,
                   download=True,
                   transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                      batch_size=1024,
                      num_workers=10,
                      shuffle=False)
```

MLP for Cifar-10

Case 2

❖ ReLU, He and Adam

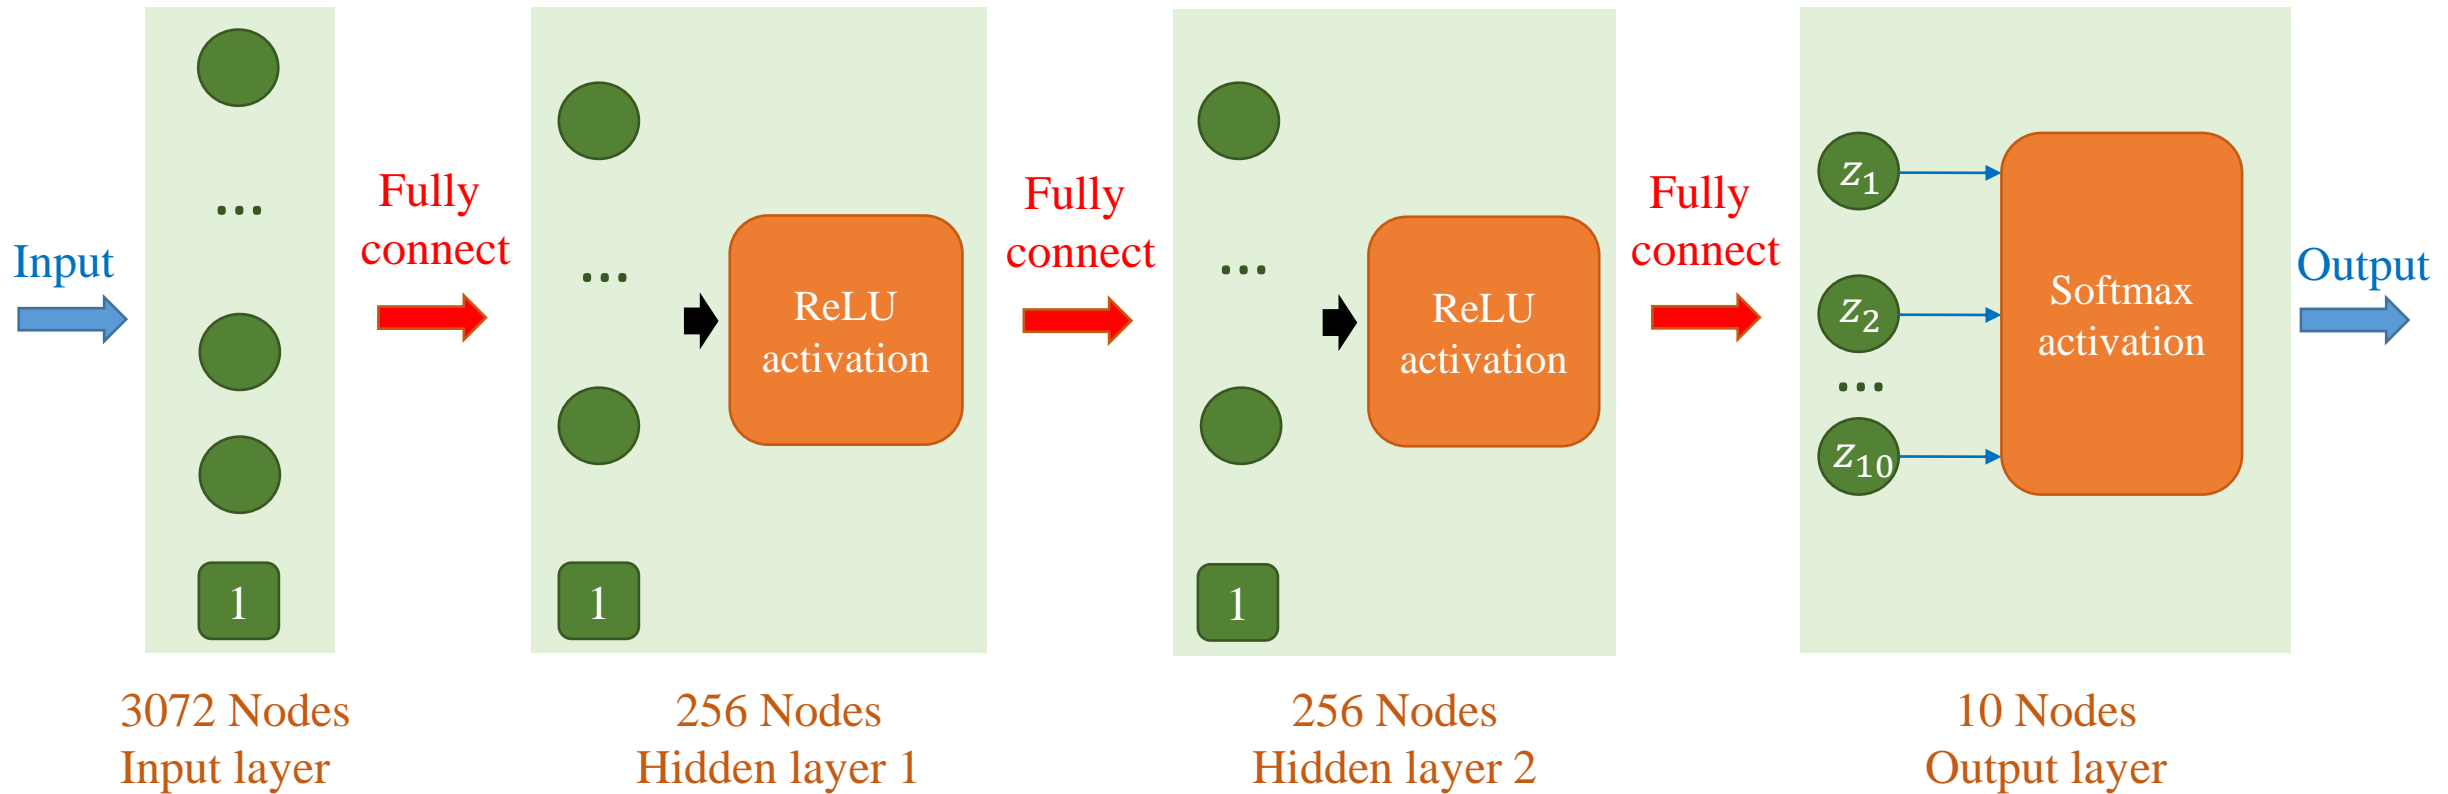


Perform disappointedly

MLP for Cifar-10

Case 3

❖ ReLU, He and Adam: add more layers



MLP for Cifar-10

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(32*32*3, 256),
    nn.ReLU(),
    nn.Linear(256, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                               nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                        lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,0.5, 0.5),
                               (0.5,0.5, 0.5))])



trainset = CIFAR10(root='data',
                   train=True,
                   download=True,
                   transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                       batch_size=1024,
                       num_workers=10,
                       shuffle=False)
```

MLP for Cifar-10

Case 3

❖ ReLU, He and Adam

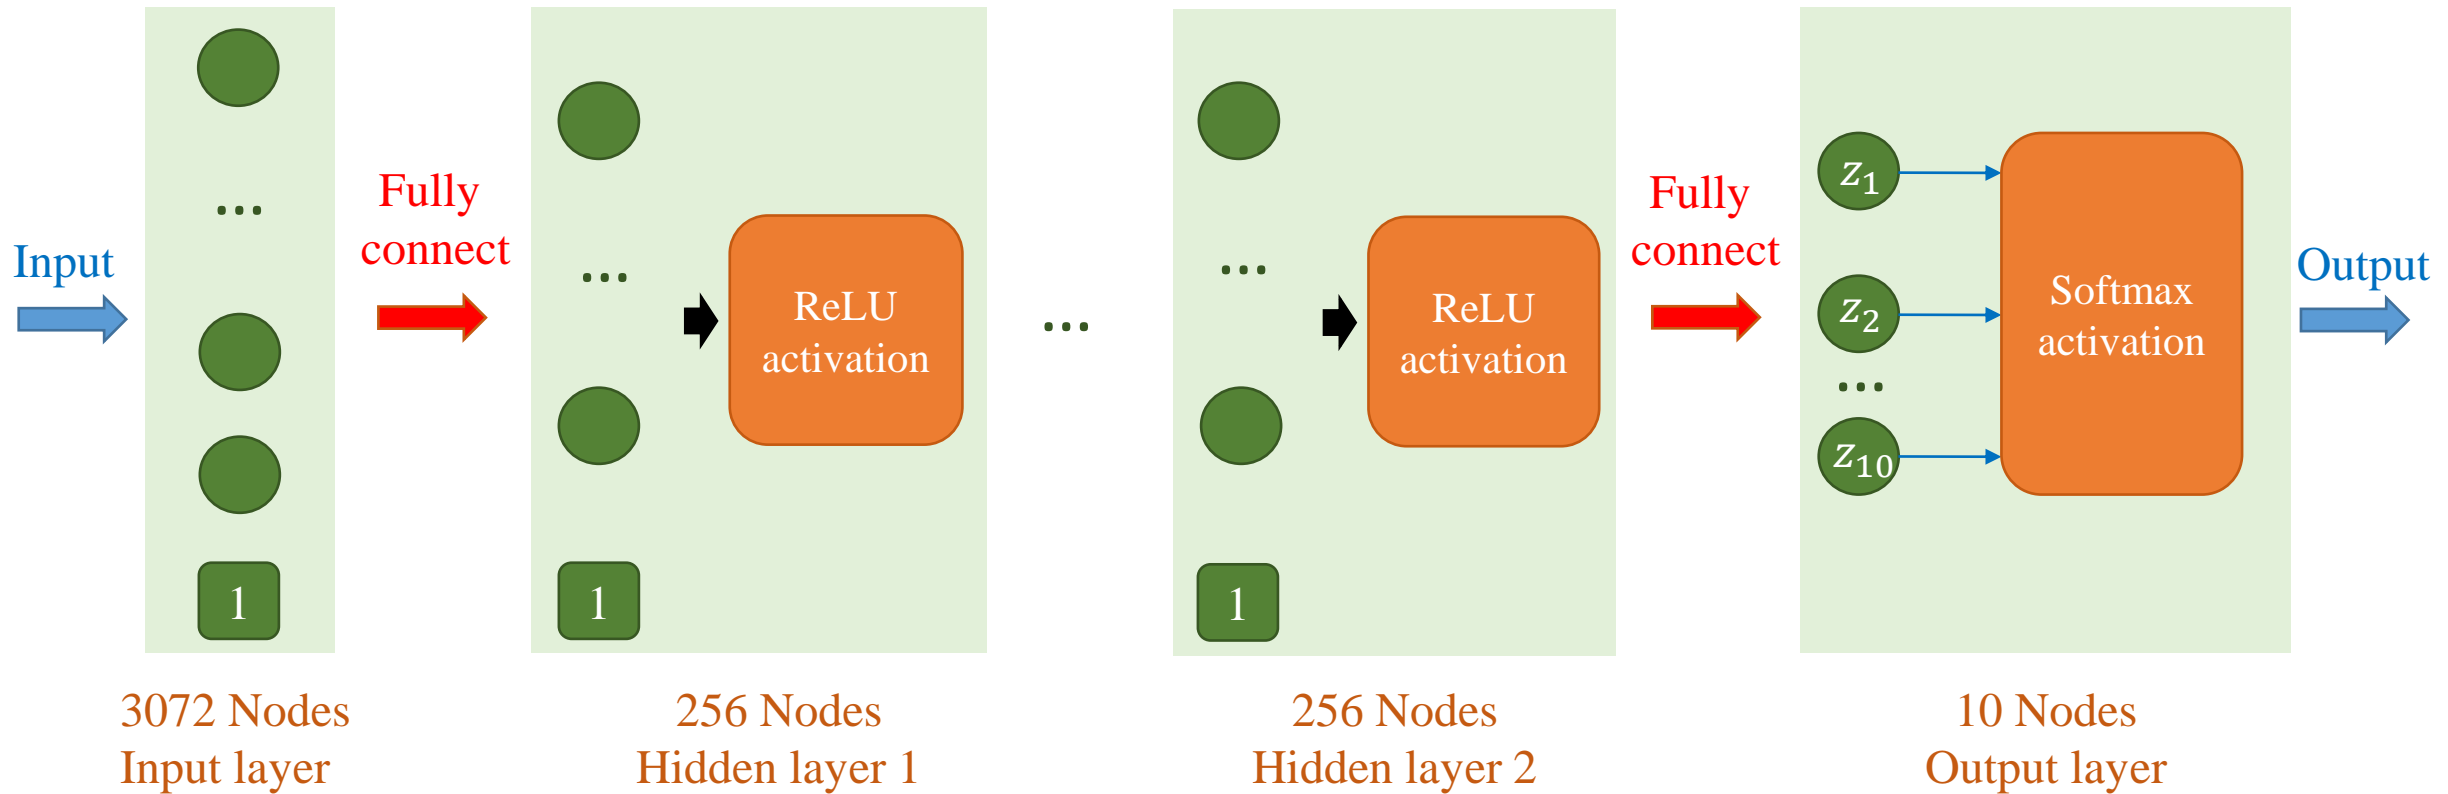


Still Perform poorly

MLP for Cifar-10

Case 4

❖ ReLU, He and Adam



MLP for Cifar-10

Case 4

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(), nn.Linear(32*32*3, 256),
    nn.ReLU(), nn.Linear(256, 256),
    nn.ReLU(), nn.Linear(256, 256),
    nn.ReLU(), nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                               nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                        lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,0.5, 0.5),
                               (0.5,0.5, 0.5))])


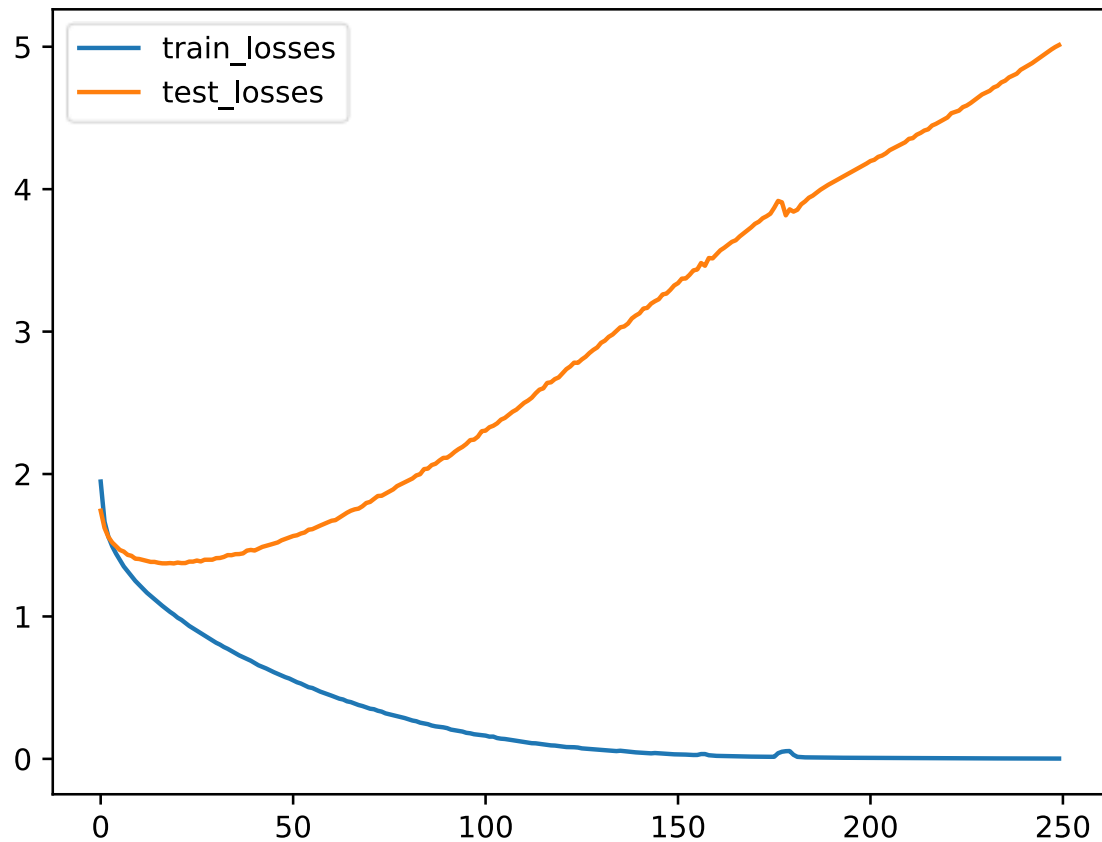
trainset = CIFAR10(root='data',
                   train=True,
                   download=True,
                   transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                 train=False,
                 download=True,
                 transform=transform)
testloader = DataLoader(testset,
                      batch_size=1024,
                      num_workers=10,
                      shuffle=False)
```


MLP for Cifar-10

Case 4

❖ ReLU, He and Adam: Using 3 hidden layers




Perform even worse

Outline

- **MLP Limitation**
- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Some Examples**
- **Application to Cifar10**

From MLP to CNN

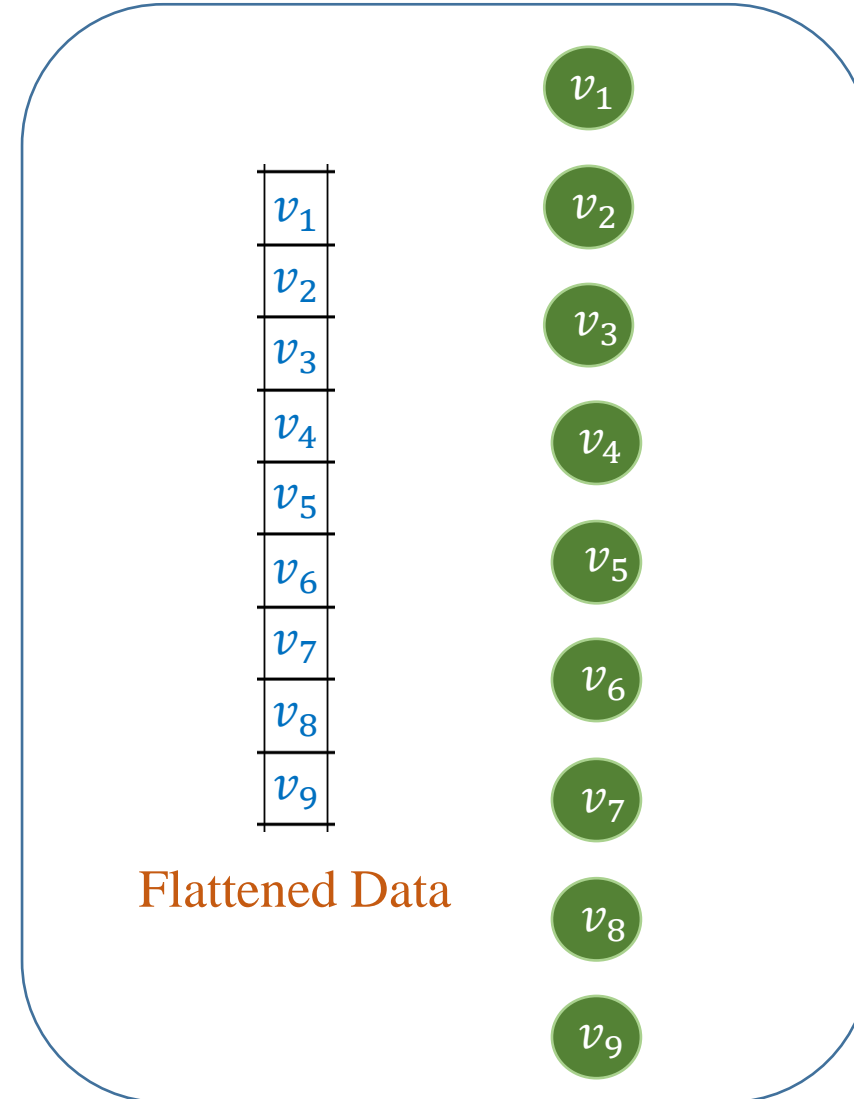
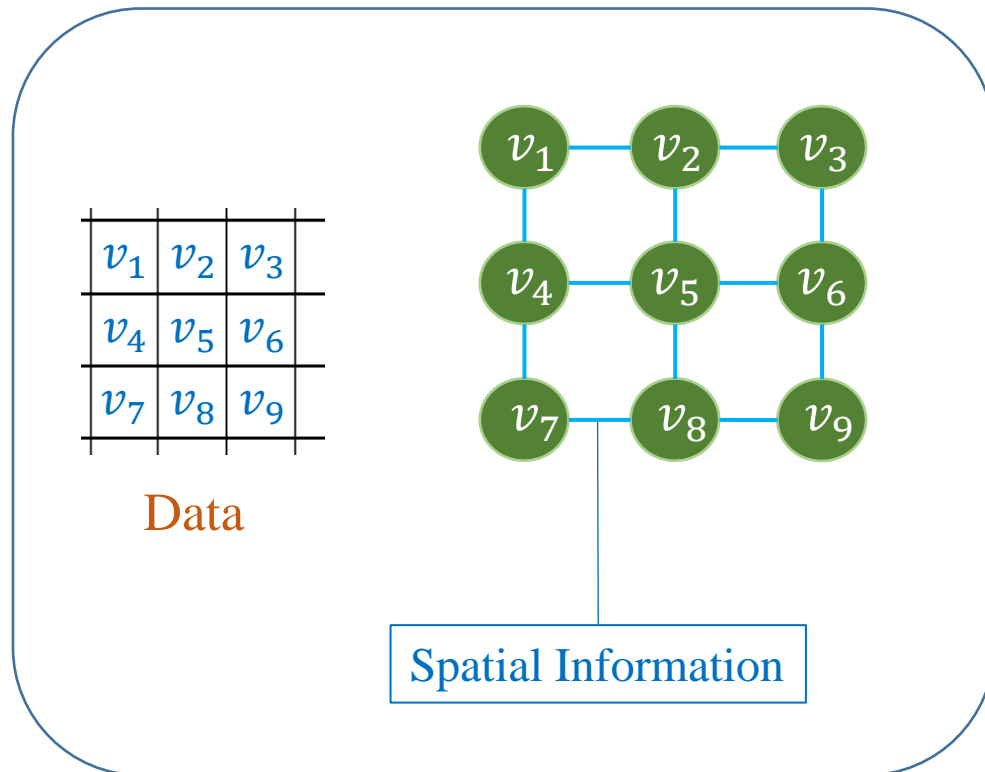
❖ Multi-layer Perceptron



Problem: Remove spatial information of the data
Inefficiently have a large amount of parameters

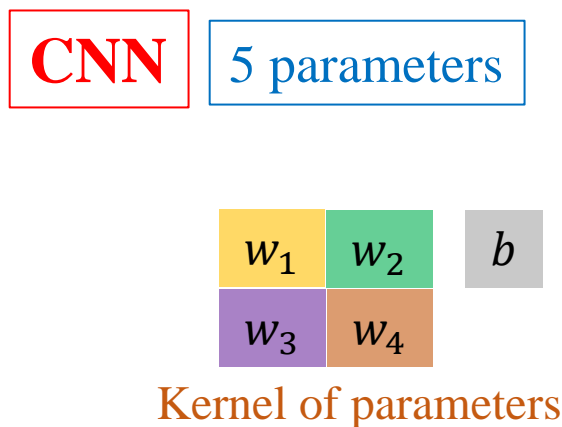

From MLP to CNN

❖ Problem of flattening data






Remove spatial information of the data

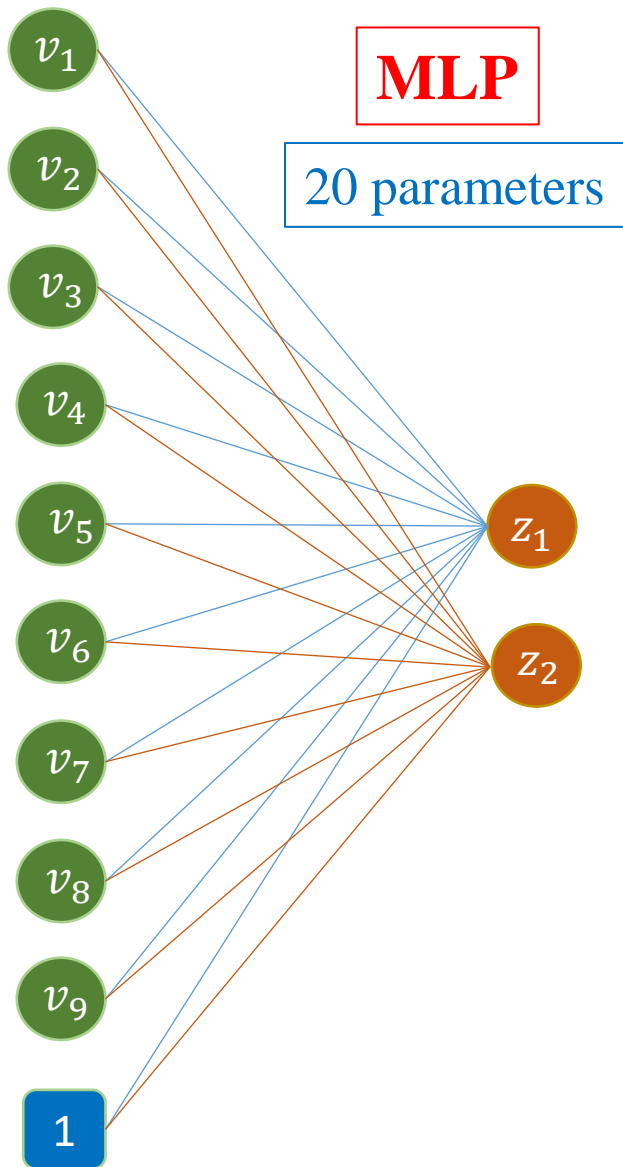
From MLP to CNN




$$m_1 = v_1 w_1 + v_2 w_2 + v_4 w_3 + v_5 w_4 + b$$







From MLP to CNN



CNN 10 parameters



Kernel 1 \neq Kernel 2




Global vs. Local?

Parameter size?

From MLP to CNN

❖ Understand convolution



Shape=(1,3,3)

(Channel=1, Height=3, Width=3)



*



Shape=(1,2,2)

#parameters (+bias) = 5

#channels of data **must** = #channels of kernel



Convolution

❖ How many cases?

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data **D**

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel **K**

Bias **b** = 0.0

<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1	<table><tr><td>0</td><td>0</td><td>1</td><td>2</td><td>2</td></tr><tr><td>1</td><td>2</td><td>2</td><td>1</td><td>2</td></tr><tr><td>0</td><td>2</td><td>0</td><td>2</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td><td>0</td><td>1</td></tr></table>	0	0	1	2	2	1	2	2	1	2	0	2	0	2	1	0	1	1	1	0	1	0	0	0	1
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									
0	0	1	2	2																																																																									
1	2	2	1	2																																																																									
0	2	0	2	1																																																																									
0	1	1	1	0																																																																									
1	0	0	0	1																																																																									

Convolution

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data **D**

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel **K**

m_1		

Output

Data size = 5×5
Kernel size = 3×3
Stride = 1

$$\begin{aligned} m_1 &= 0 \times 0.0 + 0 \times 0.1 + 1 \times -0.1 + \\ &\quad 1 \times -0.2 + 2 \times 0.0 + 2 \times 0.1 + \\ &\quad 0 \times 0.0 + 2 \times 0.0 + 0 \times 0.1 \end{aligned}$$



$$m_1 = -0.1$$

Convolution

❖ Example

$$S_o = \left\lfloor \frac{S_D - K}{S} \right\rfloor + 1$$

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data **D**

Data size = 5×5

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel **K**

Kernel size = 3×3

-0.1	-0.1	-0.2
0.3	-0.2	0.1
0.3	-0.3	0.1

Output

Stride = 1

Convolution

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data **D**

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel **K**

m_1	

Output

Data size = 5×5
Kernel size = 3×3
Stride = 2

$$\begin{aligned} m_1 &= 0 \times 0.0 + 0 \times 0.1 + 1 \times -0.1 + \\ &\quad 1 \times -0.2 + 2 \times 0.0 + 2 \times 0.1 + \\ &\quad 0 \times 0.0 + 2 \times 0.0 + 0 \times 0.1 \end{aligned}$$



$$m_1 = -0.1$$

Convolution

❖ Example

$$S_o = \left\lfloor \frac{S_D - K}{S} \right\rfloor + 1$$

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data **D**

Data size = 5×5

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel **K**

Kernel size = 3×3

-0.1	-0.2
0.3	0.1

Output

Stride = 2

Convolutional Neural Network

❖ Understand convolution



Input Data
(3,32,32)



Convolve with
1 kernel (3,5,5)




Feature map
(1,28,28)



Input Data
(3,32,32)



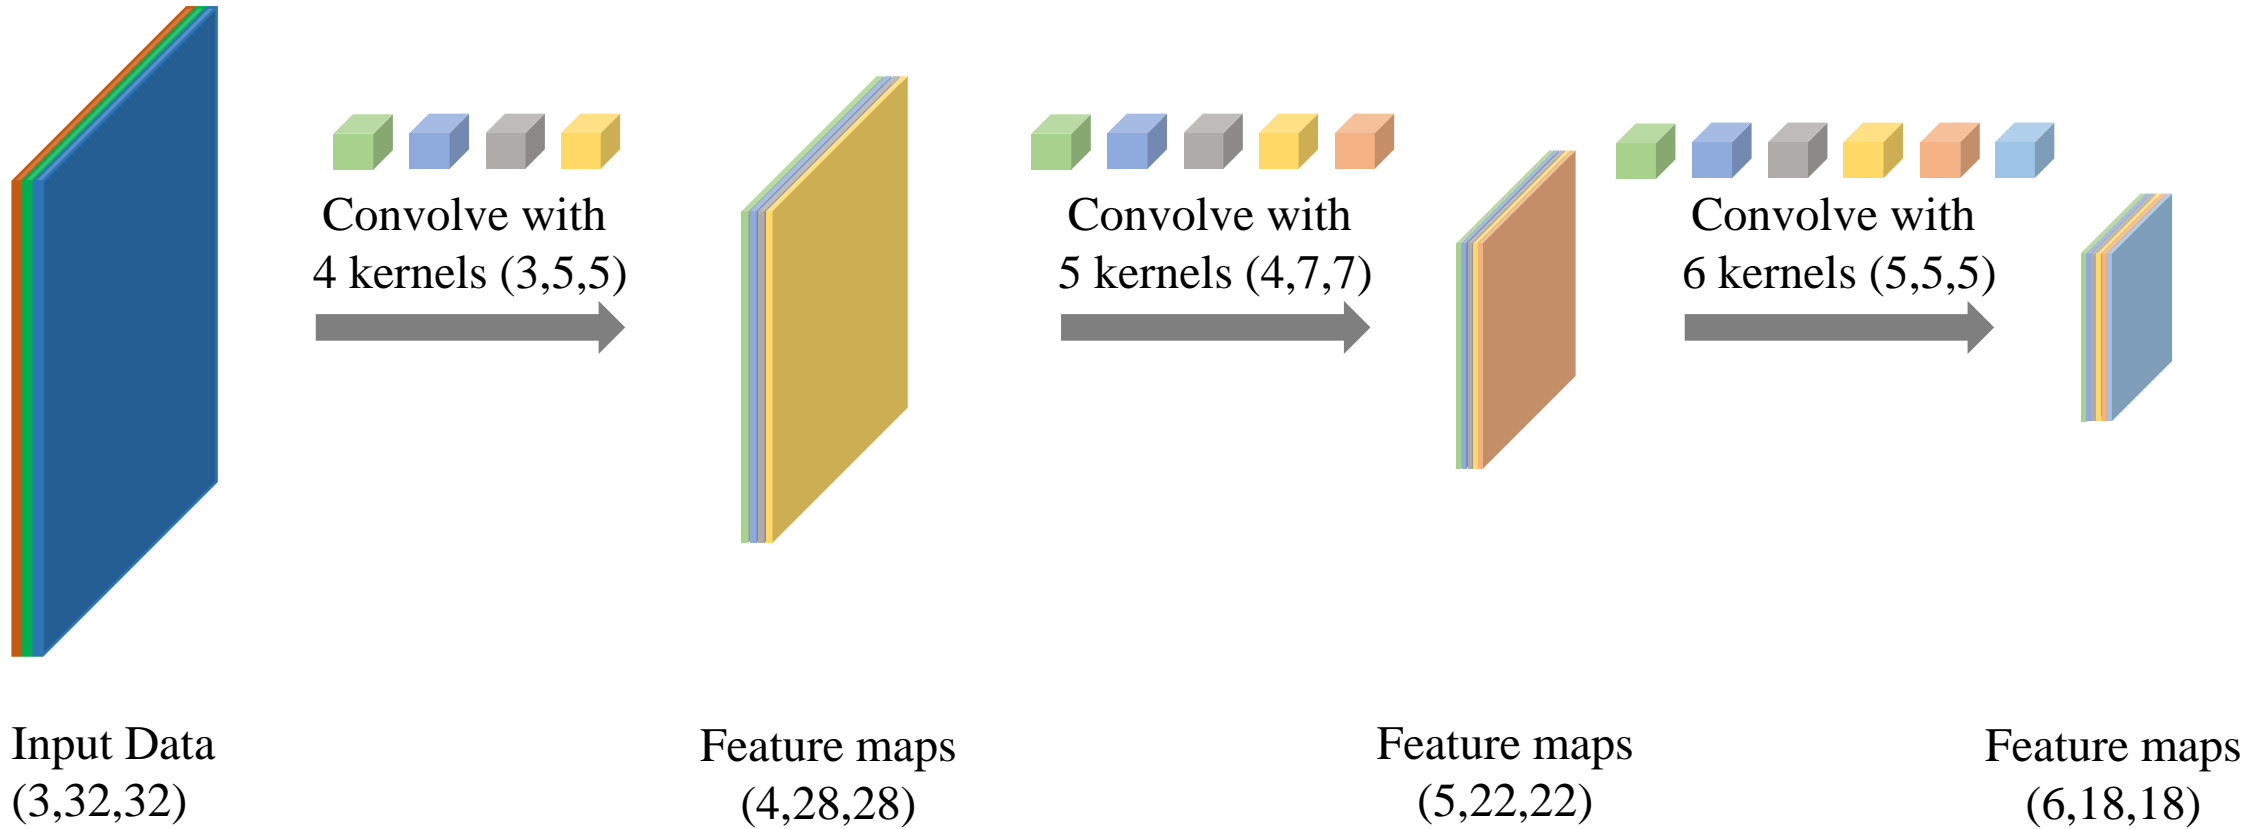
Convolve with
4 kernels (3,5,5)



Feature maps
(4,28,28)

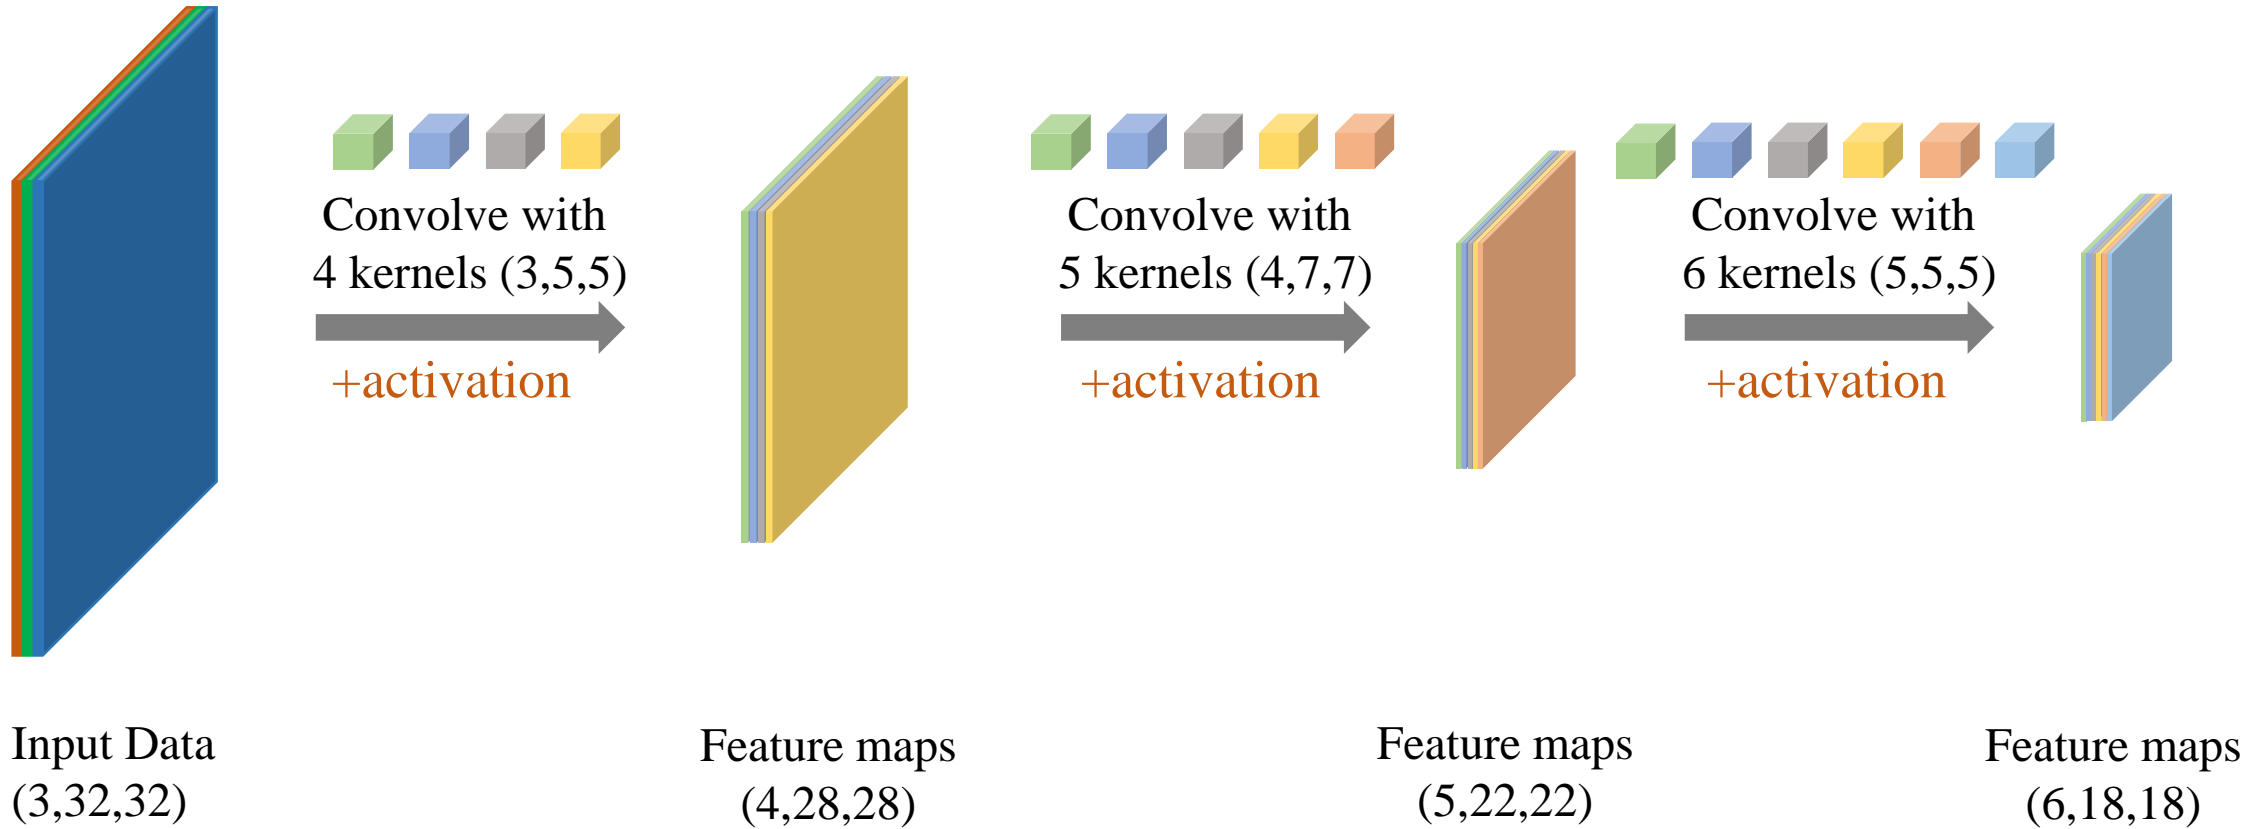
Convolutional Neural Network

❖ A stack of convolutions



Convolutional Neural Network


❖ A stack of pairs of convolution+activation




Convolutional Neural Network

❖ Convolution layer in PyTorch

```
nn.Conv2d(in_channels, out_channels, kernel_size)
```




Input Data
(1,32,32)



Feature maps
(4,28,28)



Input Data
(3,32,32)



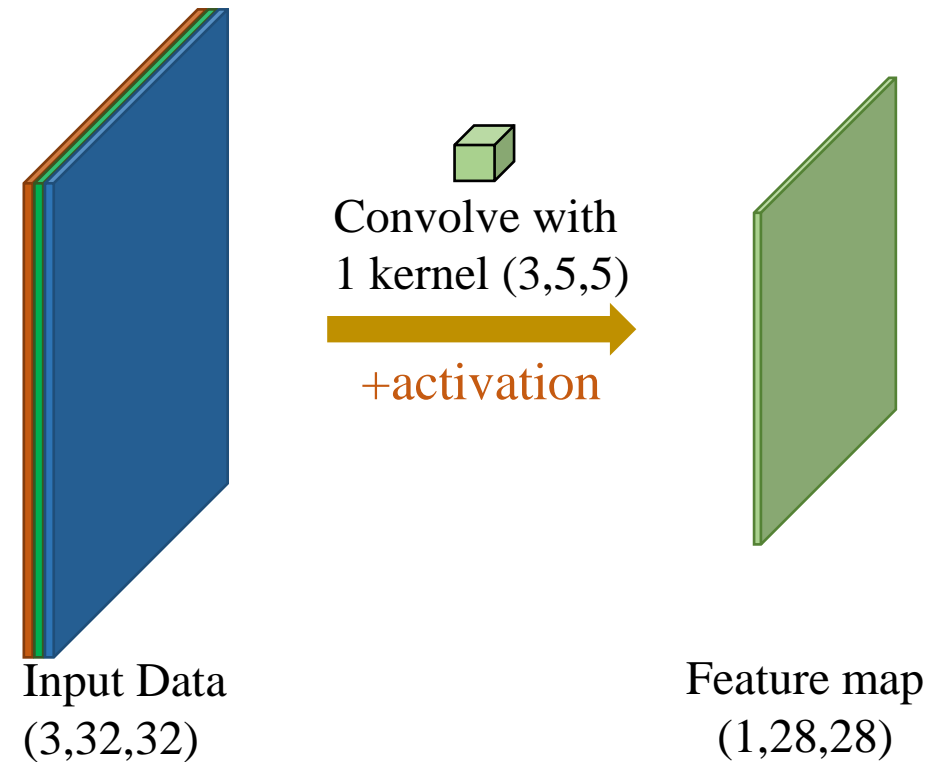
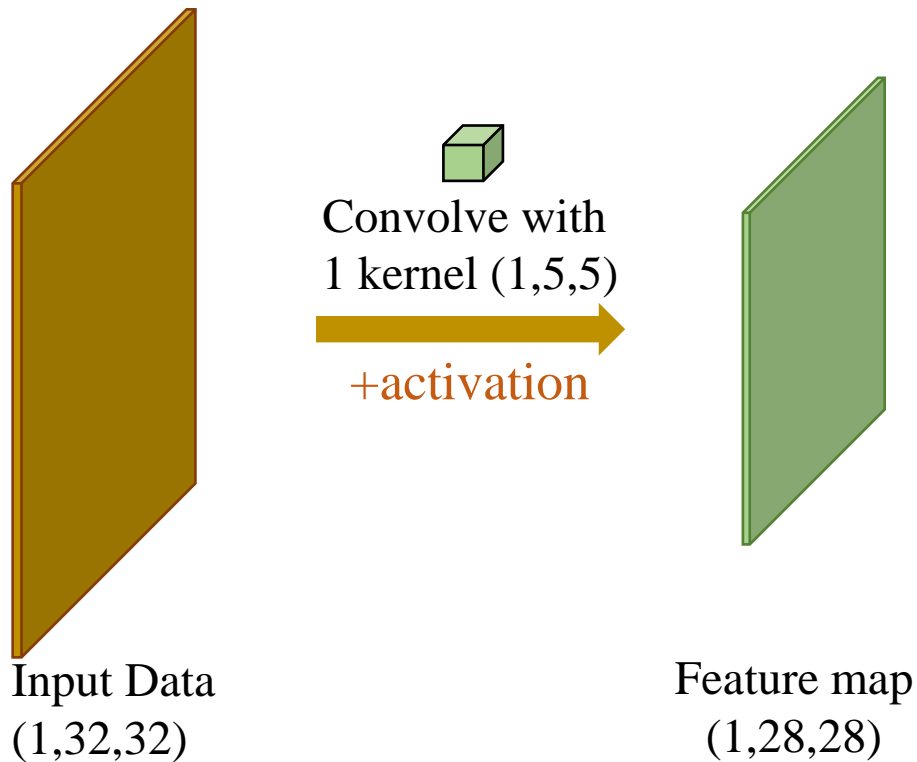
Feature maps
(4,28,28)

Convolutional Neural Network

❖ Convolution layer in PyTorch

demo

```
nn.Conv2d(in_channels, out_channels, kernel_size)  
nn.ReLU()
```



Convolutional Neural Network

Fashion-MNIST dataset


Grayscale images

Resolution=28x28


Training set: 60000 samples

Testing set: 10000 samples


T-shirt



Trouser




Pullover




Dress




Coat




Sandal



Shirt



Sneaker



Bag




Ankle
Boot



Convolutional Neural Network


❖ Apply for Fashion-MNIST dataset



Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset

demo



Simple Convolutional Neural Network

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(4*4*256, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.dense(x)
        return x


model = CustomModel()
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 22, 22]	1,600
ReLU-2	[-1, 32, 22, 22]	0
Conv2d-3	[-1, 64, 16, 16]	100,416
ReLU-4	[-1, 64, 16, 16]	0
Conv2d-5	[-1, 128, 10, 10]	401,536
ReLU-6	[-1, 128, 10, 10]	0
Conv2d-7	[-1, 256, 4, 4]	1,605,888
ReLU-8	[-1, 256, 4, 4]	0
Flatten-9	[-1, 4096]	0
Linear-10	[-1, 128]	524,416
ReLU-11	[-1, 128]	0
Linear-12	[-1, 10]	1,290
Total params: 2,635,146		
Trainable params: 2,635,146		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.78		
Params size (MB): 10.05		
Estimated Total Size (MB): 10.83		

Convolutional Neural Network

❖ Apply for Cifar-10 dataset

demo

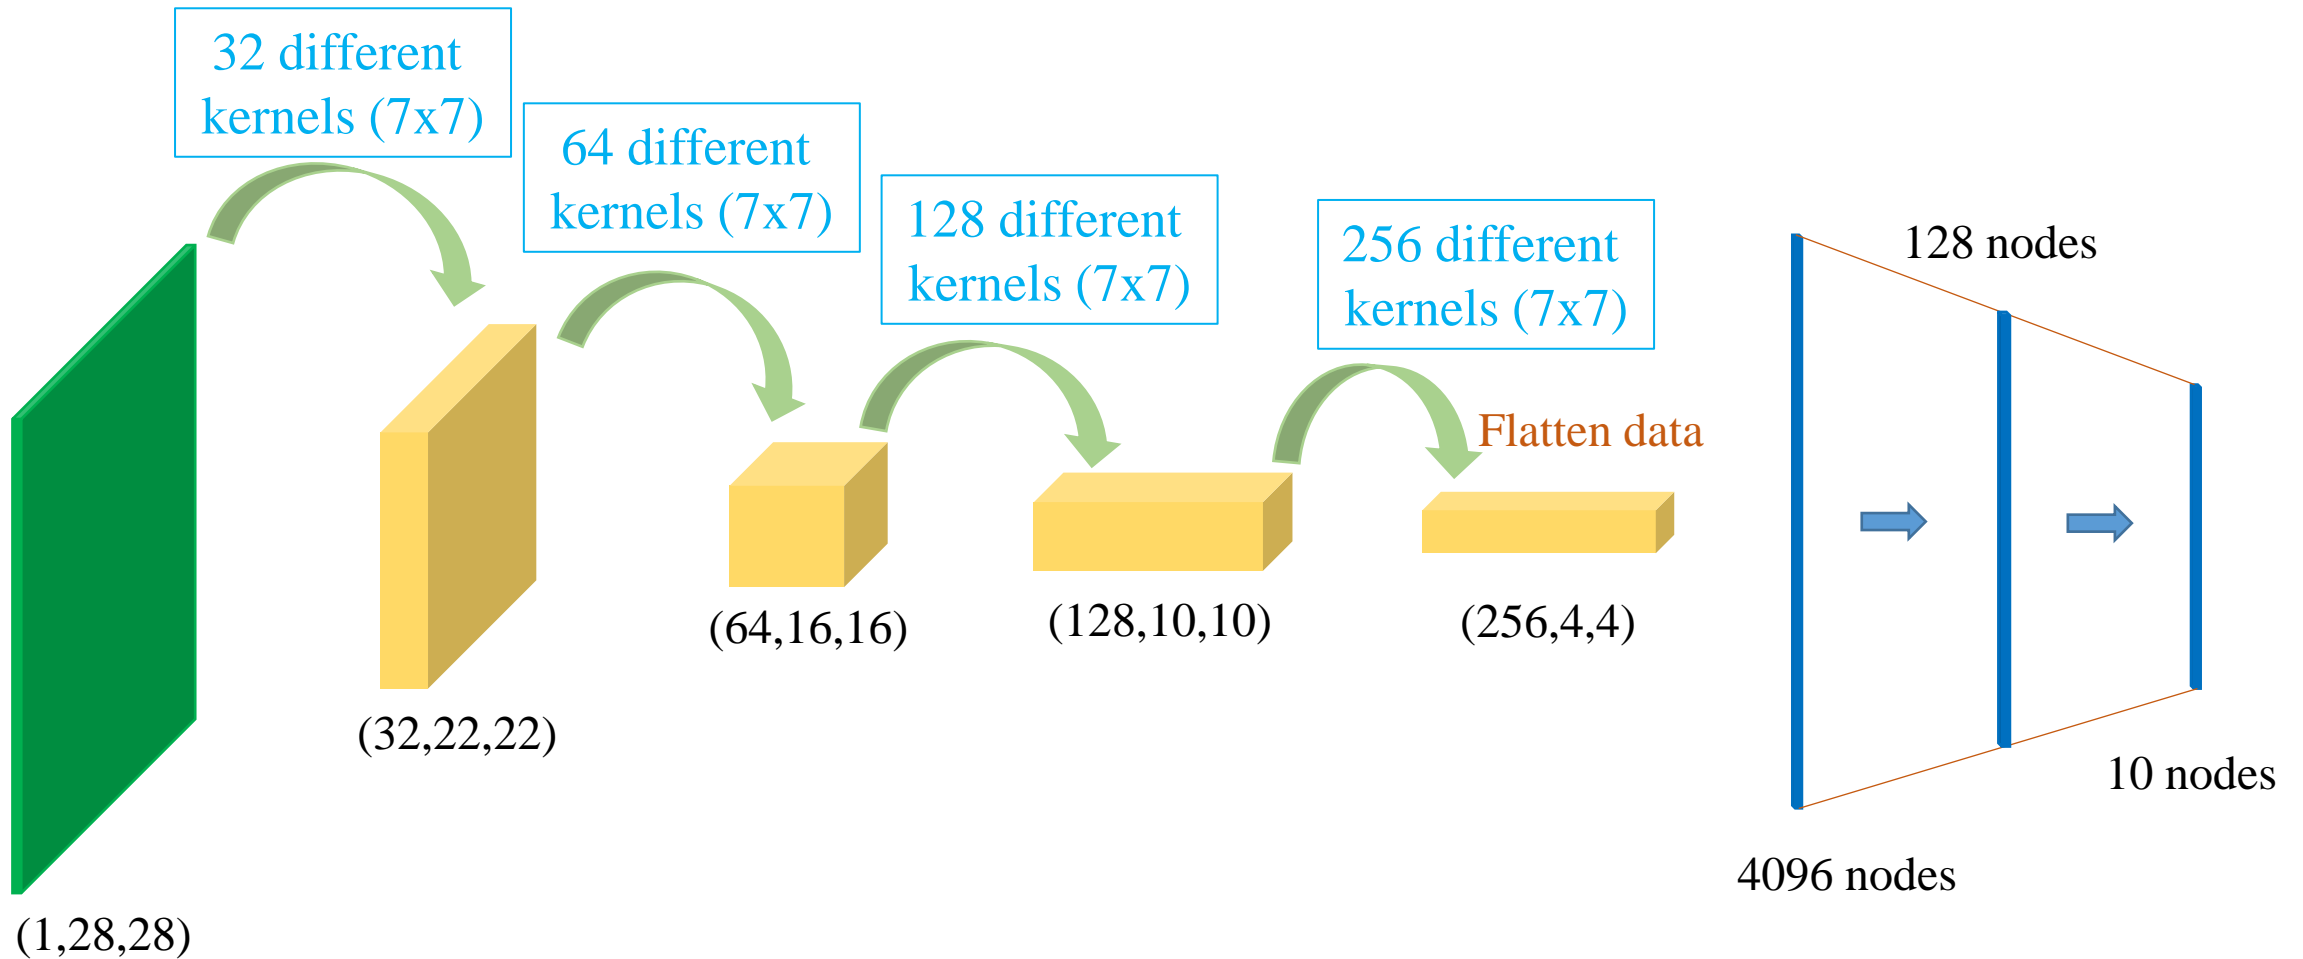


Outline

- **MLP Limitation**
- **From MLP to CNN**
- **Feature Map Down-sampling**
- **Some Examples**
- **Application to Cifar10**

Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset: case 1



```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(4*4*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)

```

```

# Load FashionMNIST dataset
transform = Compose([ToTensor(),
                      Normalize((0.5,),
                                (0.5,))])

trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)

trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = FashionMNIST(root='data',
                      train=False,
                      download=True,
                      transform=transform)



testloader = DataLoader(testset,
                      batch_size=1024,
                      num_workers=10,
                      shuffle=False)

# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)

```

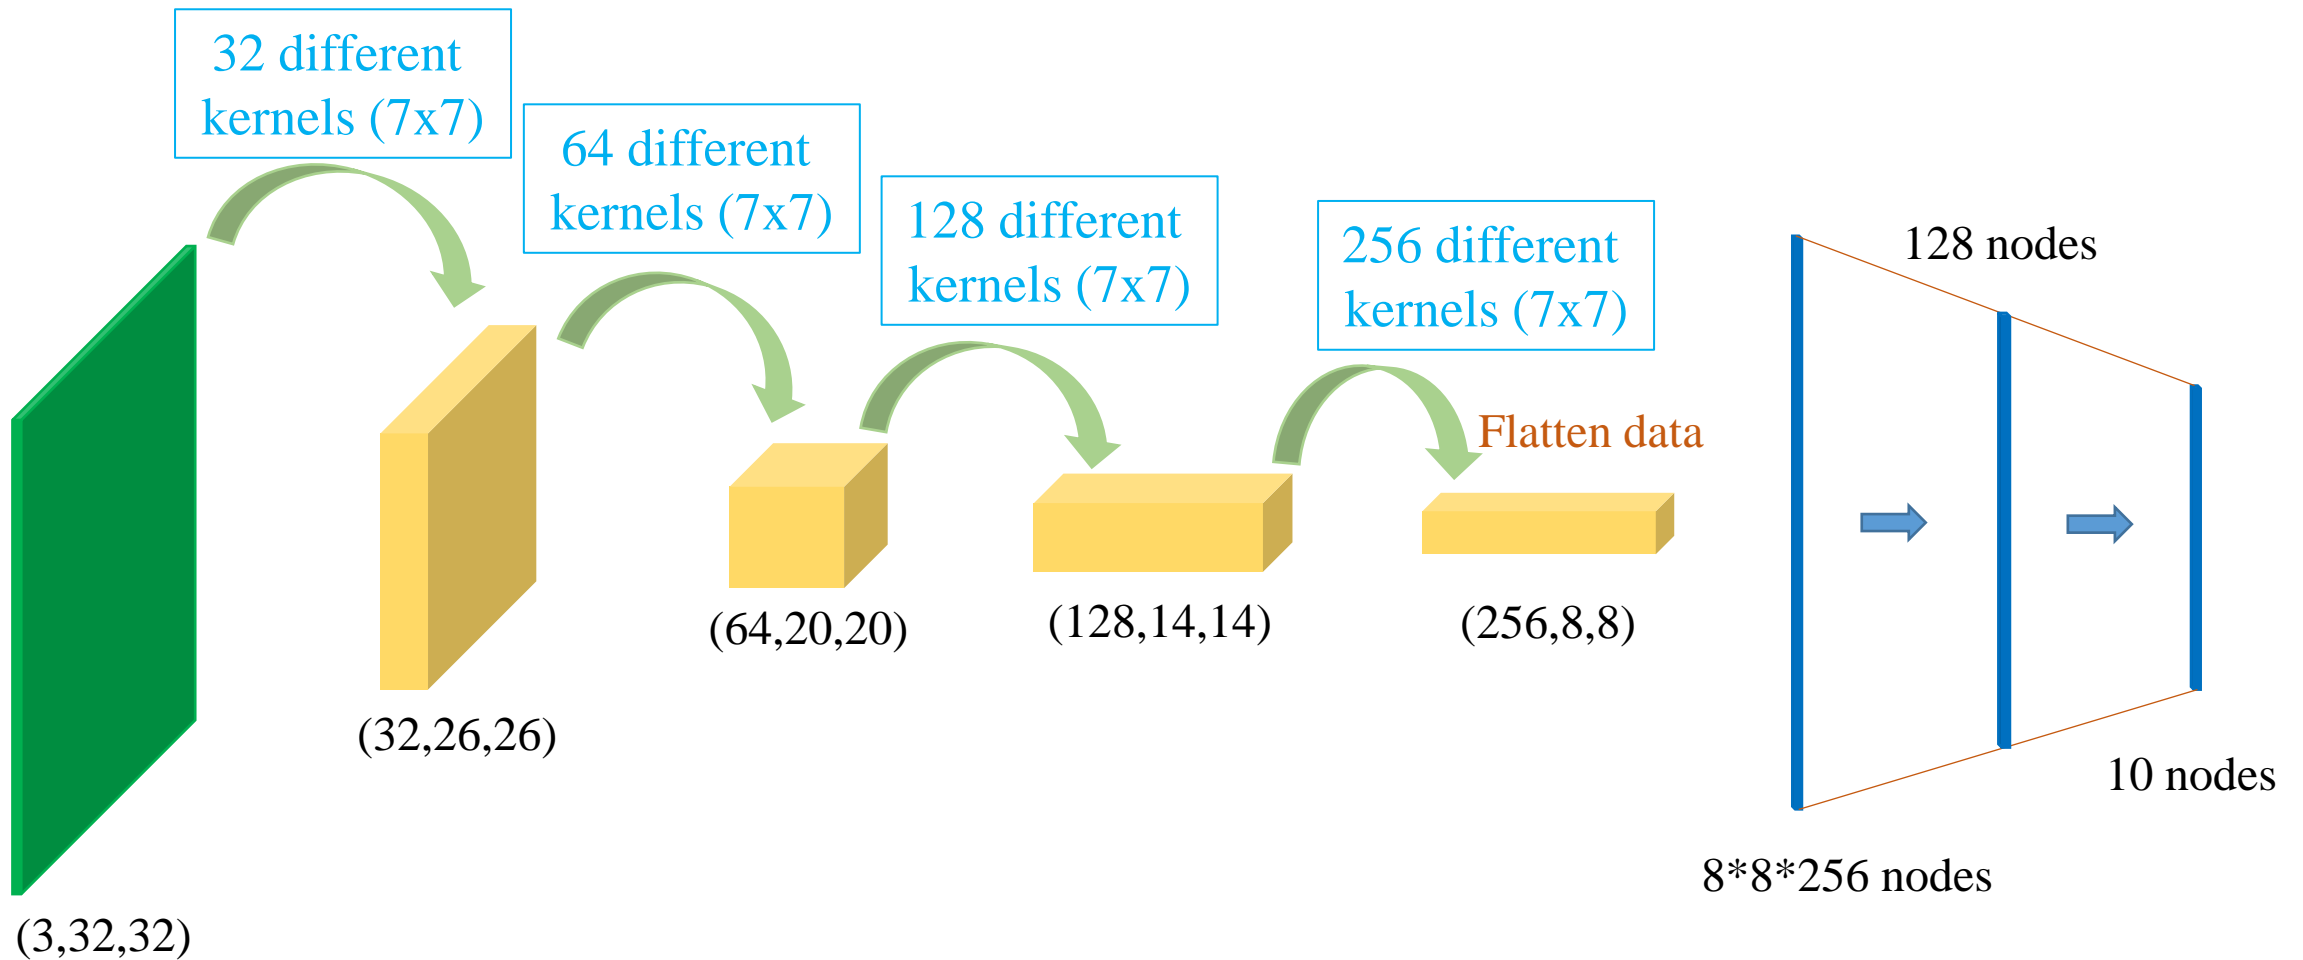
Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset: case 1



Convolutional Neural Network

❖ Apply for Cifar-10 dataset: case 2



```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(8*8*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)

```

```

# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                      Normalize((0.5,0.5, 0.5),
                                (0.5,0.5, 0.5))])

trainset = CIFAR10(root='data',
                    train=True,
                    download=True,
                    transform=transform)
trainloader = DataLoader(trainset,
                          batch_size=1024,
                          num_workers=10,
                          shuffle=True,
                          drop_last=True)



testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                         batch_size=1024,
                         num_workers=10,
                         shuffle=False)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)

```

Convolutional Neural Network

❖ Apply for Cifar-10 dataset: case 2



Test Accuracy from MLP: ~53%

Further Reading

❖ Reading

<https://cs231n.github.io/convolutional-networks/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

