

**AI VIETNAM**  
**All-in-One Course**  
**(TA Session)**

# **Image Retrieval**

## **Project**



**AI VIET NAM**  
@aivietnam.edu.vn

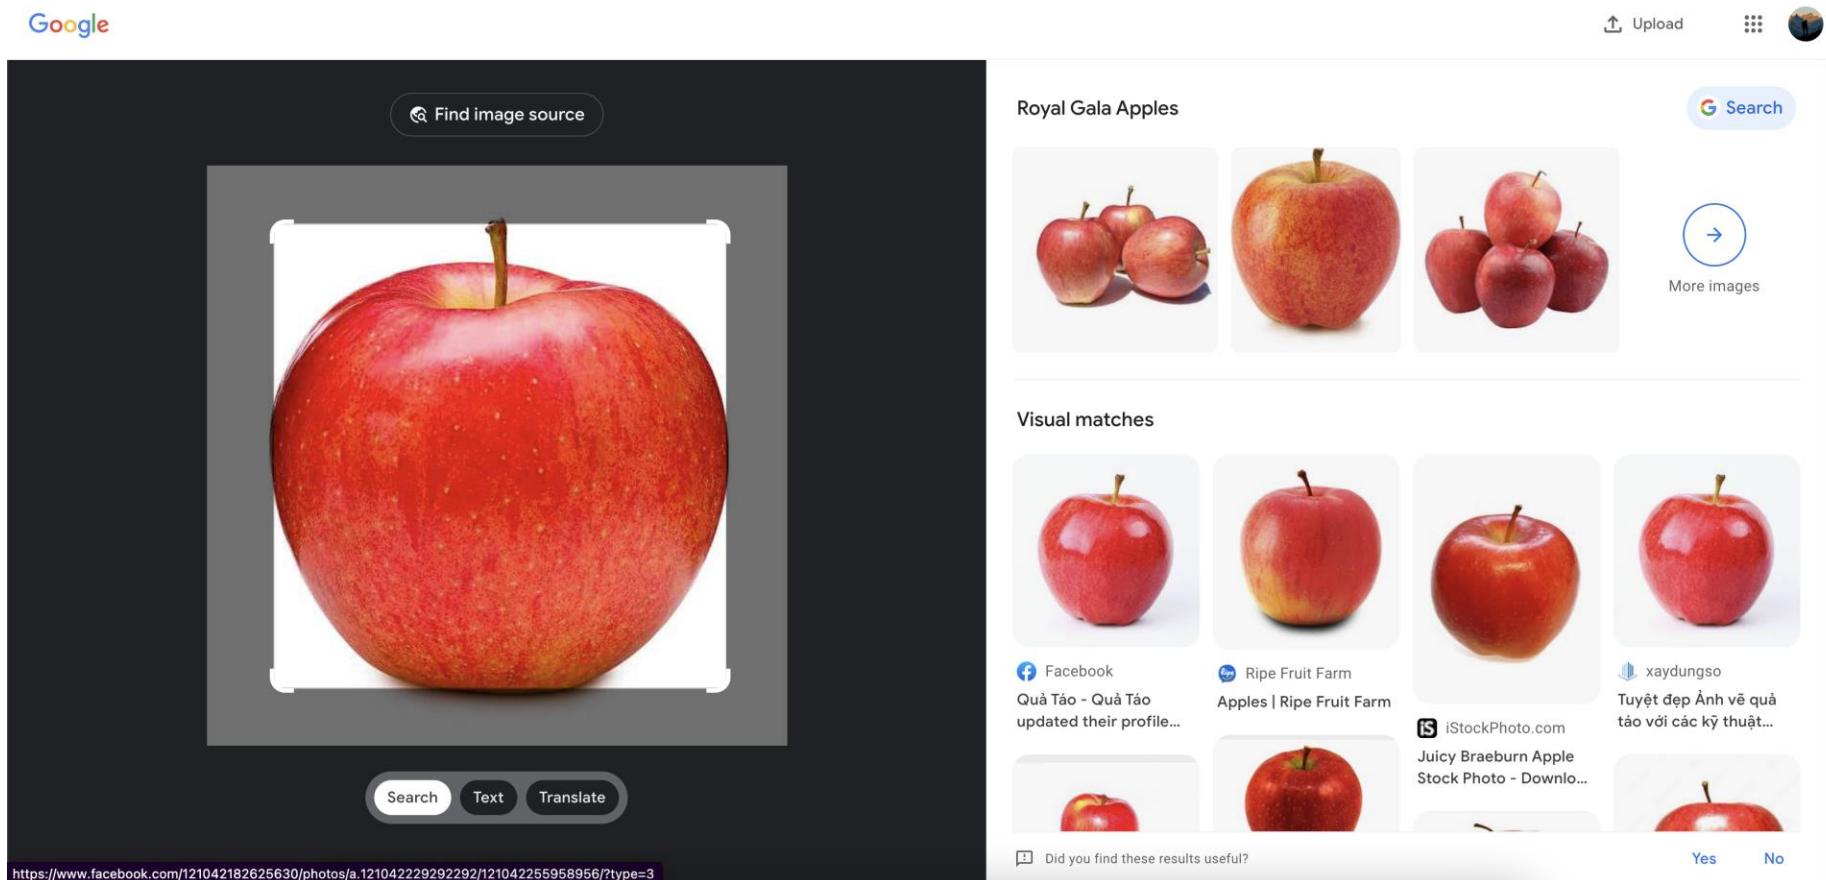
**Dinh-Thang Duong – TA**

# Outline

- Introduction
- Baseline
- Apply pre-trained ViT
- Question

# Introduction

## ❖ Getting Started



# Introduction

## ❖ Getting Started

### Input Image

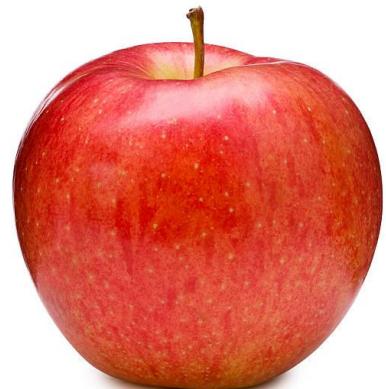


Image Search Engine

Google images

Search any image with Google Lens

Drag an image here or [upload a file](#)

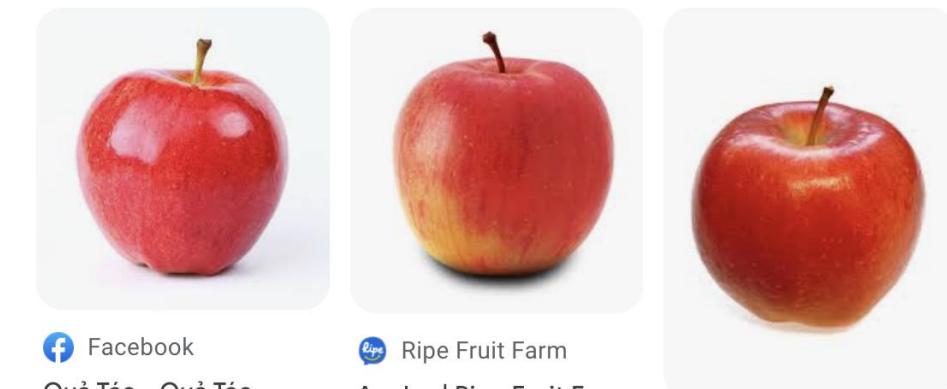
OR

Paste image link

Search

### Results

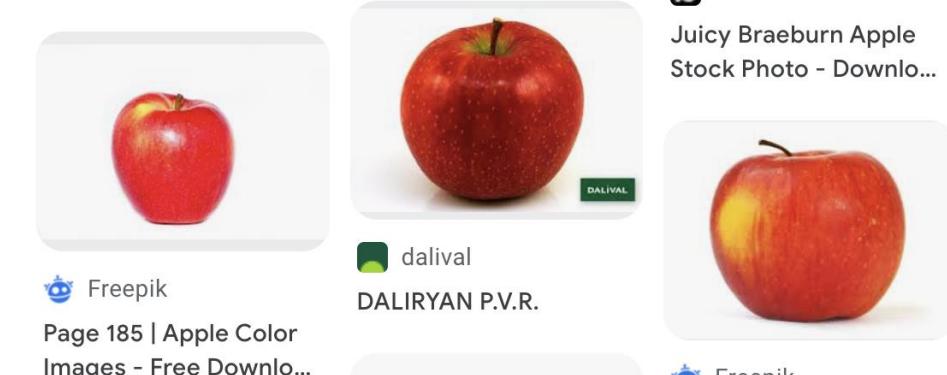
Visual matches



Facebook  
Quả Táo - Quả Táo updated their profile...

Ripe Fruit Farm  
Apples | Ripe Fruit Farm

iStockPhoto.com  
Juicy Braeburn Apple Stock Photo - Download



FreePik  
Page 185 | Apple Color Images - Free Download

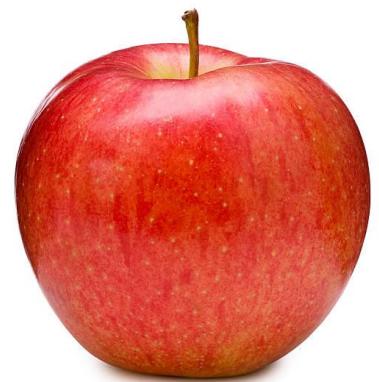
DALIRYAN P.V.R.



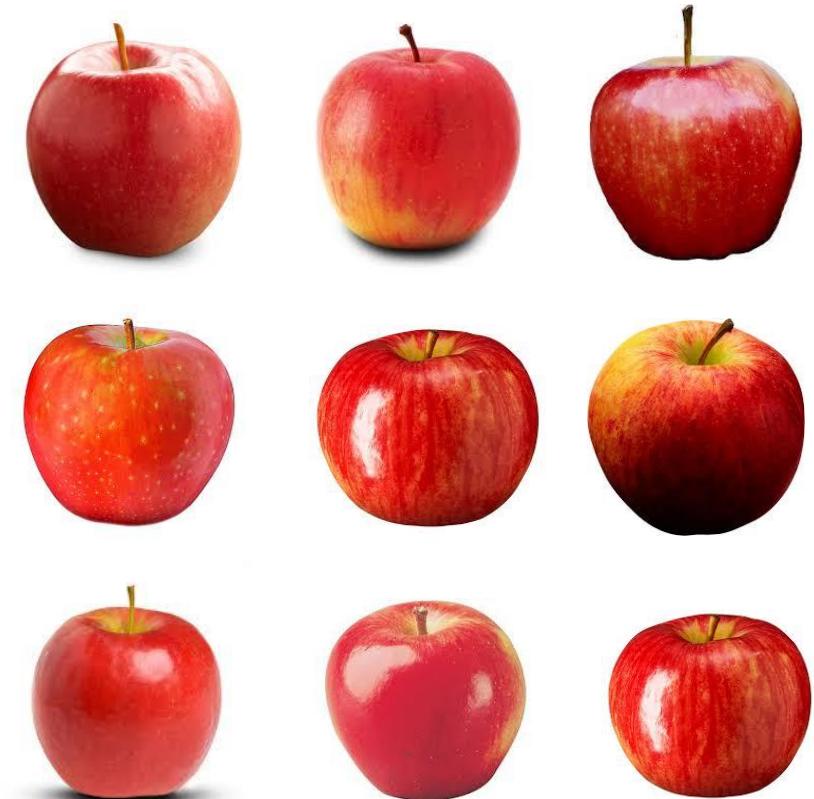
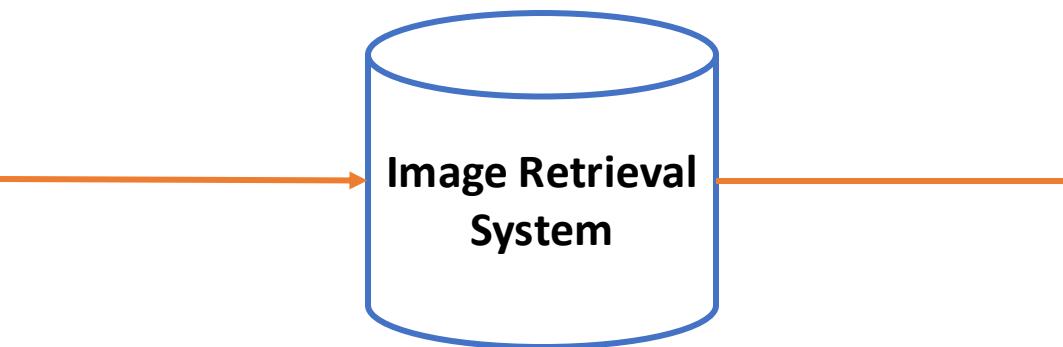
FreePik  
Page 11 | Redyellow Images - Free Download

# Introduction

## ❖ Image Retrieval



Query Image



Top K most similar images

An **image retrieval** system is a computer system used for browsing, searching and retrieving images from a large database of digital images.

# Introduction

---

## ❖ Applications



Google



Baidu



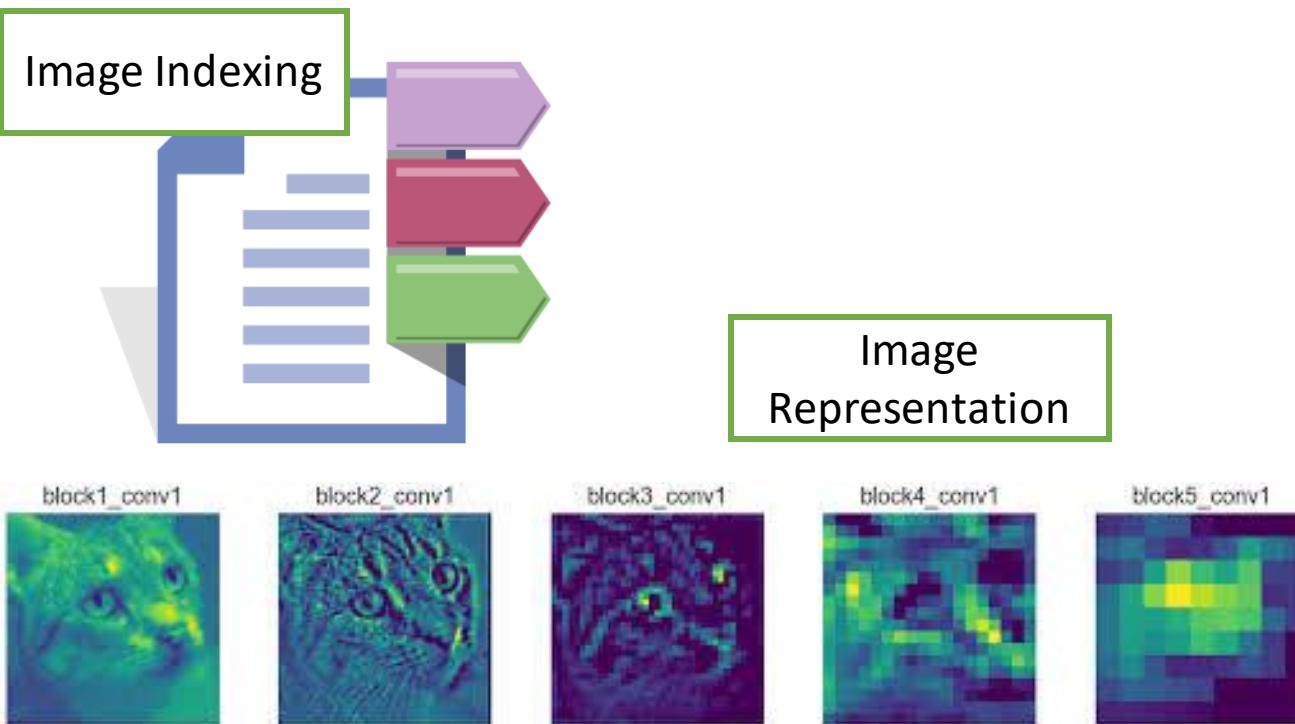
Yandex



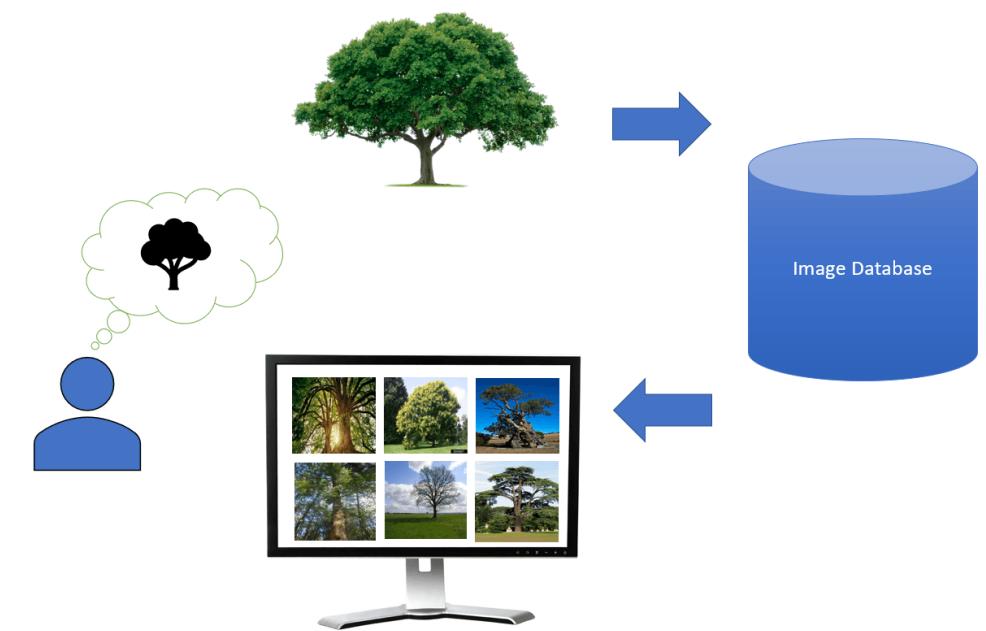
Bing

# Introduction

## ❖ Challenges



Share the same some IR challenges



How to satisfy **information need?**

# Introduction

## ❖ Description

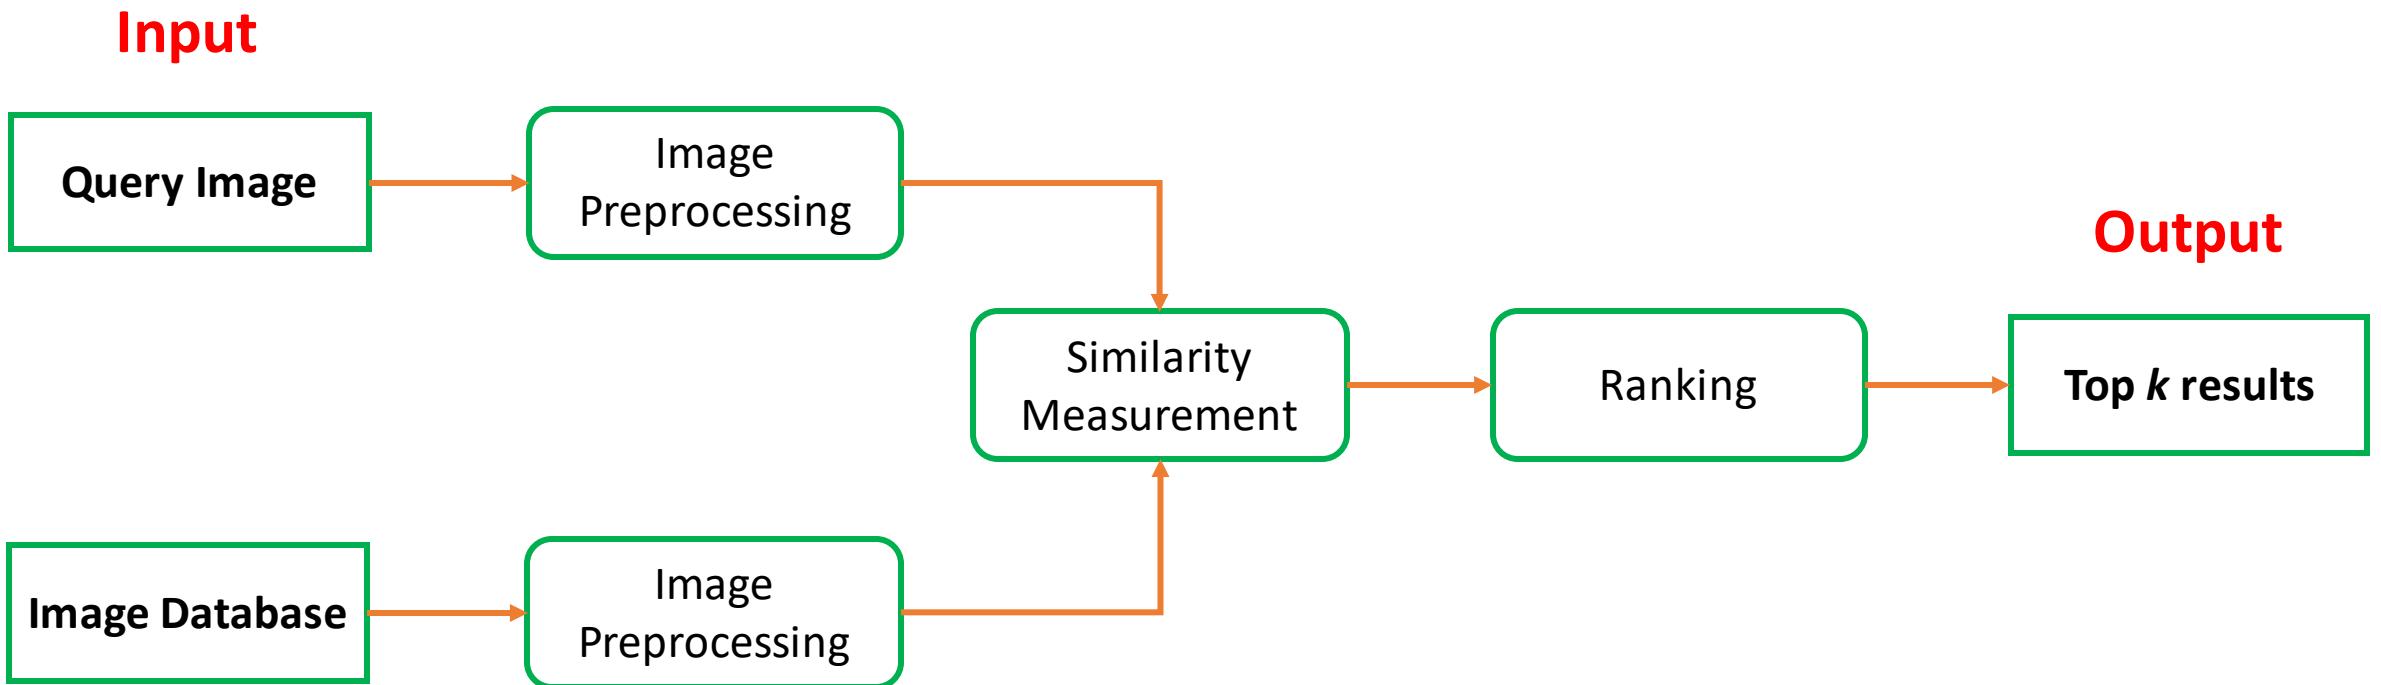
**Project Statement:** Given a list of images, build an image retrieval system that finds relevant or similar images among the source images compared to the query image.



[Download the dataset here](#)

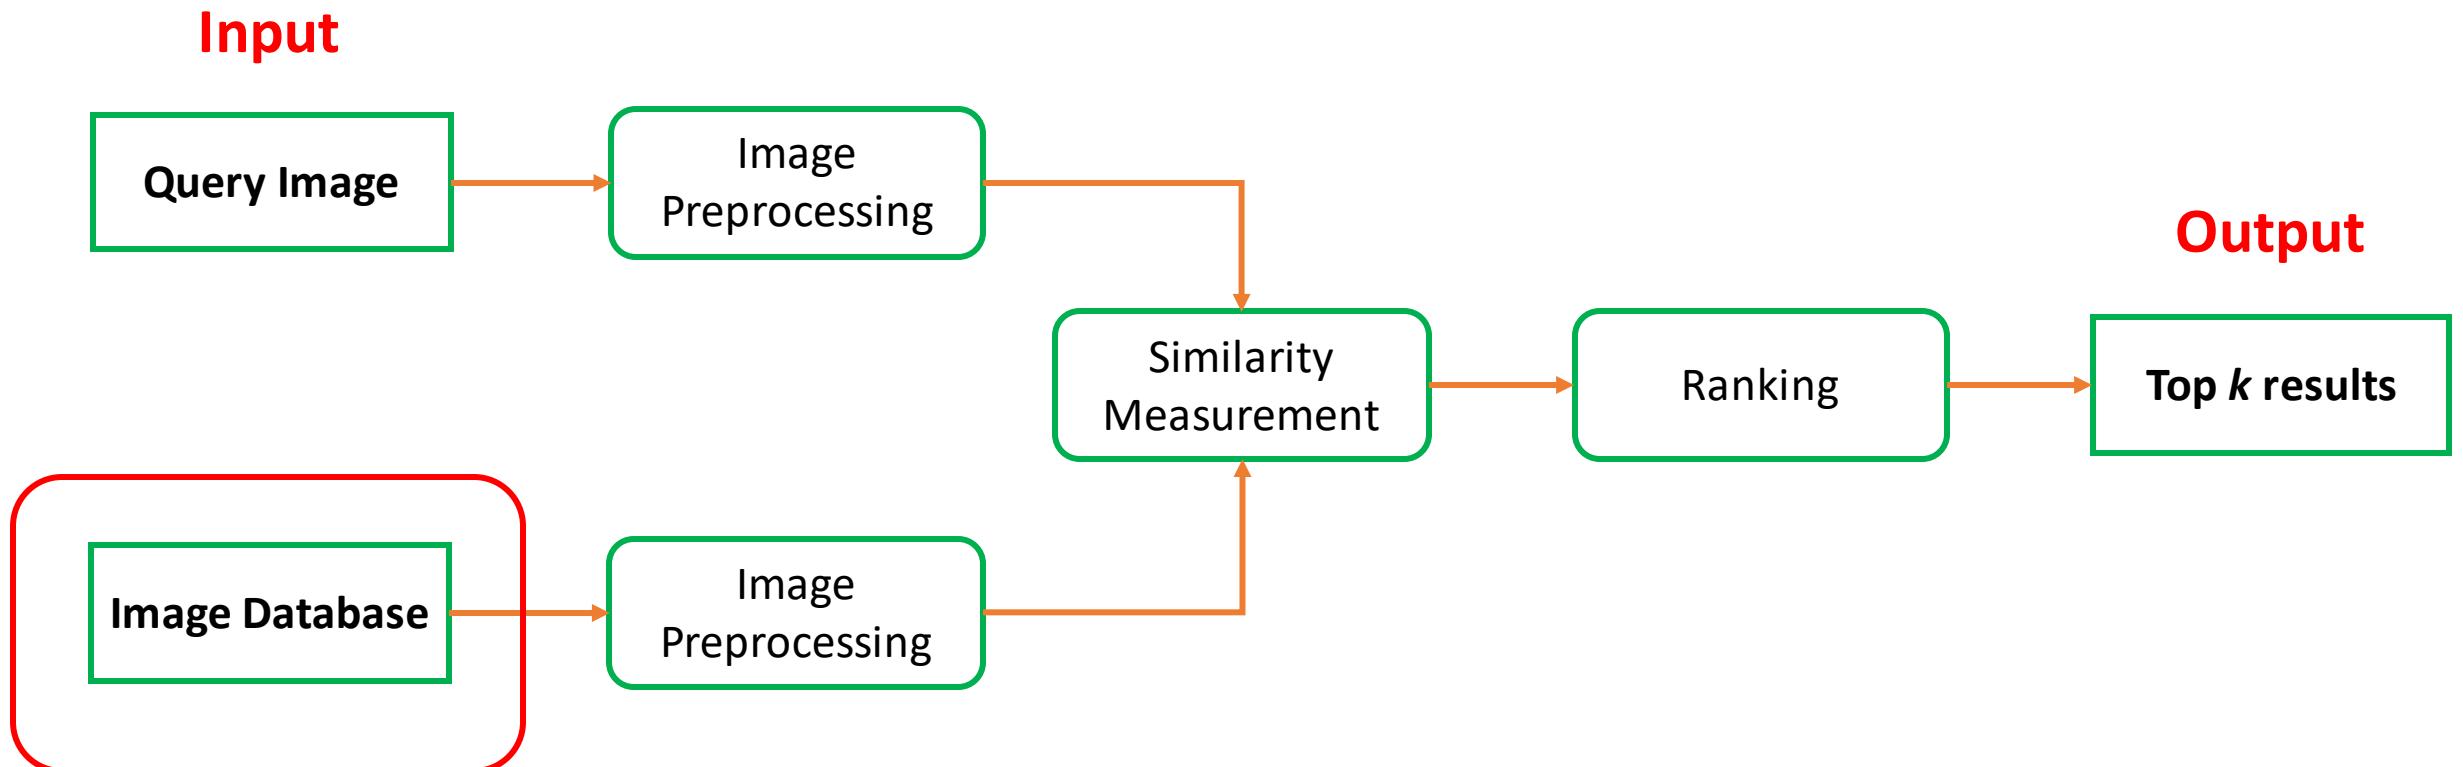
# Baseline

## ❖ Project Pipeline



# Baseline

## ❖ Project Pipeline



# Baseline

## ❖ Step 1: Import libraries and get filenames

### 1. Import necessary libraries



```
1 import os
2 import numpy as np
3 import cv2
4 import matplotlib.pyplot as plt
```

### 2. Get image filenames



```
1 # Declare source images directory path
2 dataset_dir = 'images_mr'
3
4 # Get all filenames inside "dataset_dir"
5 image_filenames = os.listdir(dataset_dir)
```

```
os.listdir(path='.'
```

Return a list containing the names of the entries in the directory given by *path*. The list is in arbitrary order, and does not include the special entries `'.'` and `'..'` even if they are present in the directory. If a file is removed from or added to the directory during the call of this function, whether a name for that file be included is unspecified.

```
1 image_filenames
```

```
'9385.jpg',
'5052.jpg',
'6271.jpg',
'5362.jpg',
'5852.jpg',
'6164.jpg',
'8962.jpg',
'3608.jpg',
'6438.jpg',
'6670.jpg',
'5018.jpg',
'8697.jpg',
```

```
1 print(
2     'Number of files: ',
3     len(image_filenames)
4 )
```

Number of files: 9908

# Baseline

## ❖ Step 2: Read images



```
1 # Declare empty list
2 src_images = []
3
4 for filename in image_filenames:
5     # Create filepath of current image
6     filepath = os.path.join(
7         dataset_dir,
8         filename
9     )
10
11    image = cv2.imread(filepath) # Read image
12    image = cv2.cvtColor(
13        image,
14        cv2.COLOR_BGR2RGB
15    ) # Convert to RGB
16    src_images.append(image)
```

**os.path.join(path, \*paths)**  
Join one or more path segments intelligently.

### Example:

- dataset\_dir = “image\_sr”
  - filename = “01.jpg”
- => filepath = “image\_sr/01.jpg”

# Baseline

## ❖ Step 2: Read images



```
1 # Declare empty list
2 src_images = []
3
4 for filename in image_filenames:
5     # Create filepath of current image
6     filepath = os.path.join(
7         dataset_dir,
8         filename
9     )
10
11     image = cv2.imread(filepath) # Read image
12     image = cv2.cvtColor(
13         image,
14         cv2.COLOR_BGR2RGB
15     ) # Convert to RGB
16     src_images.append(image)
```

`cv2.imread(path, flag)`

Loads an image from a specified file (channels stored in B G R order for color images).

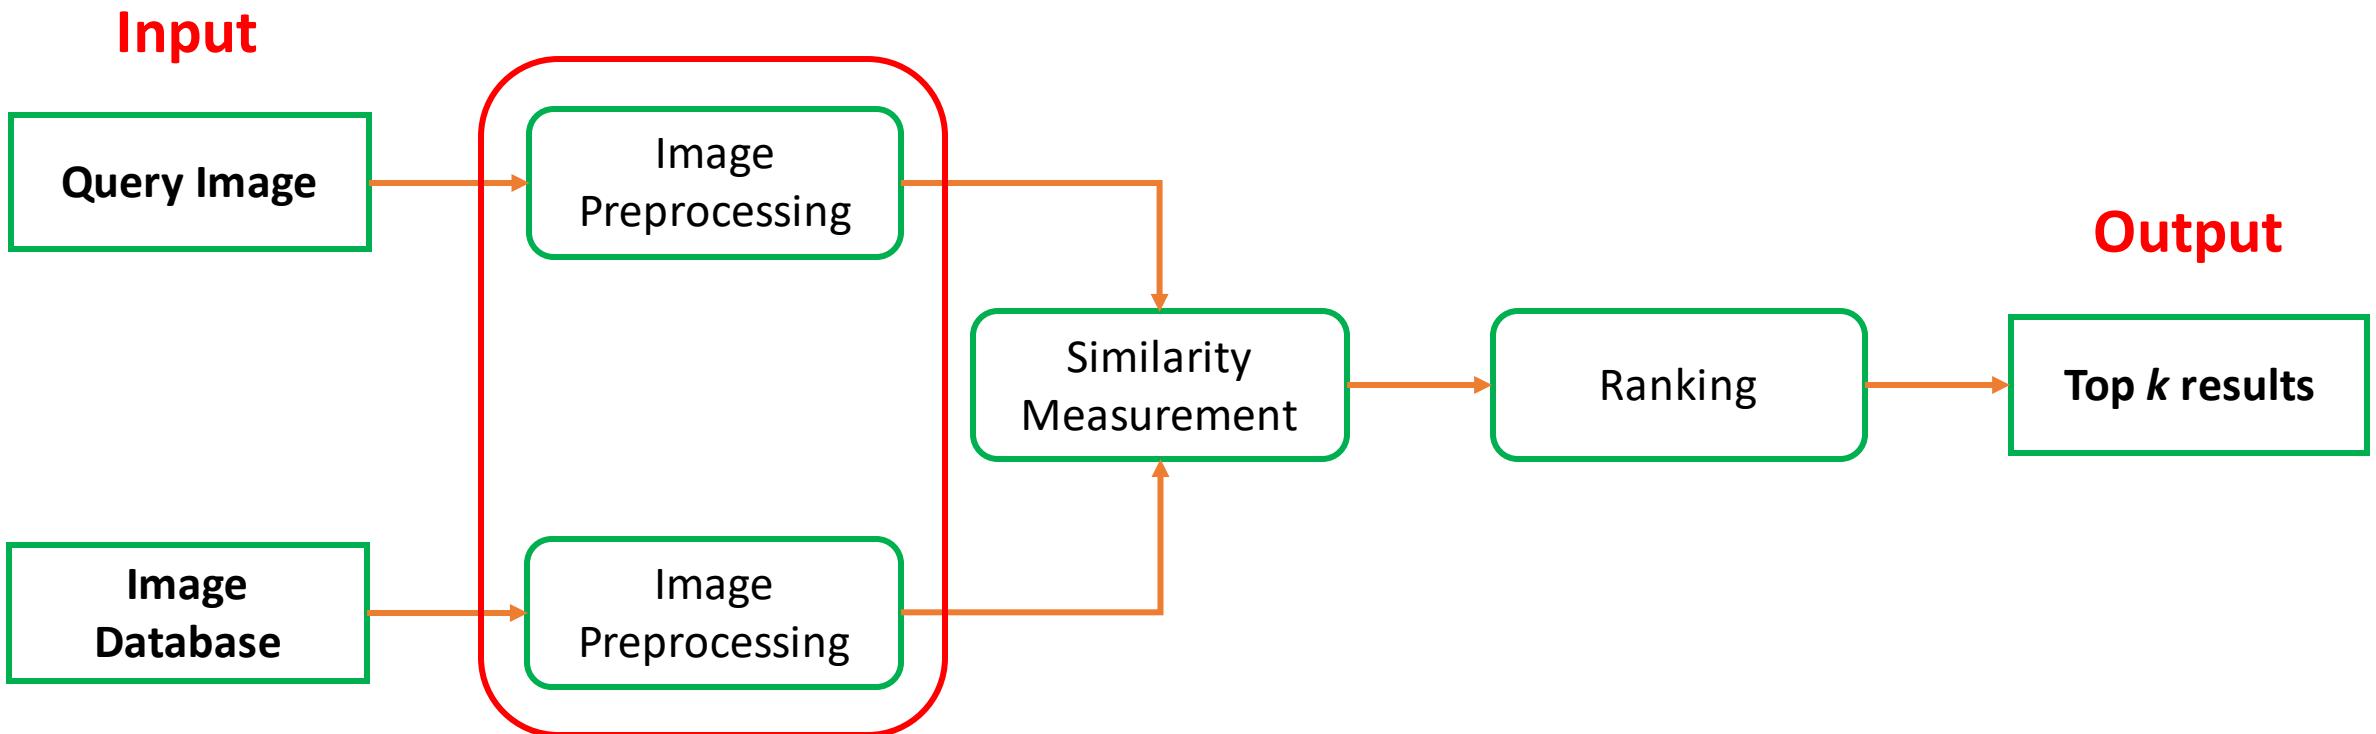
`cv2.cvtColor(src, code)`

Convert an input image from one color space to another.



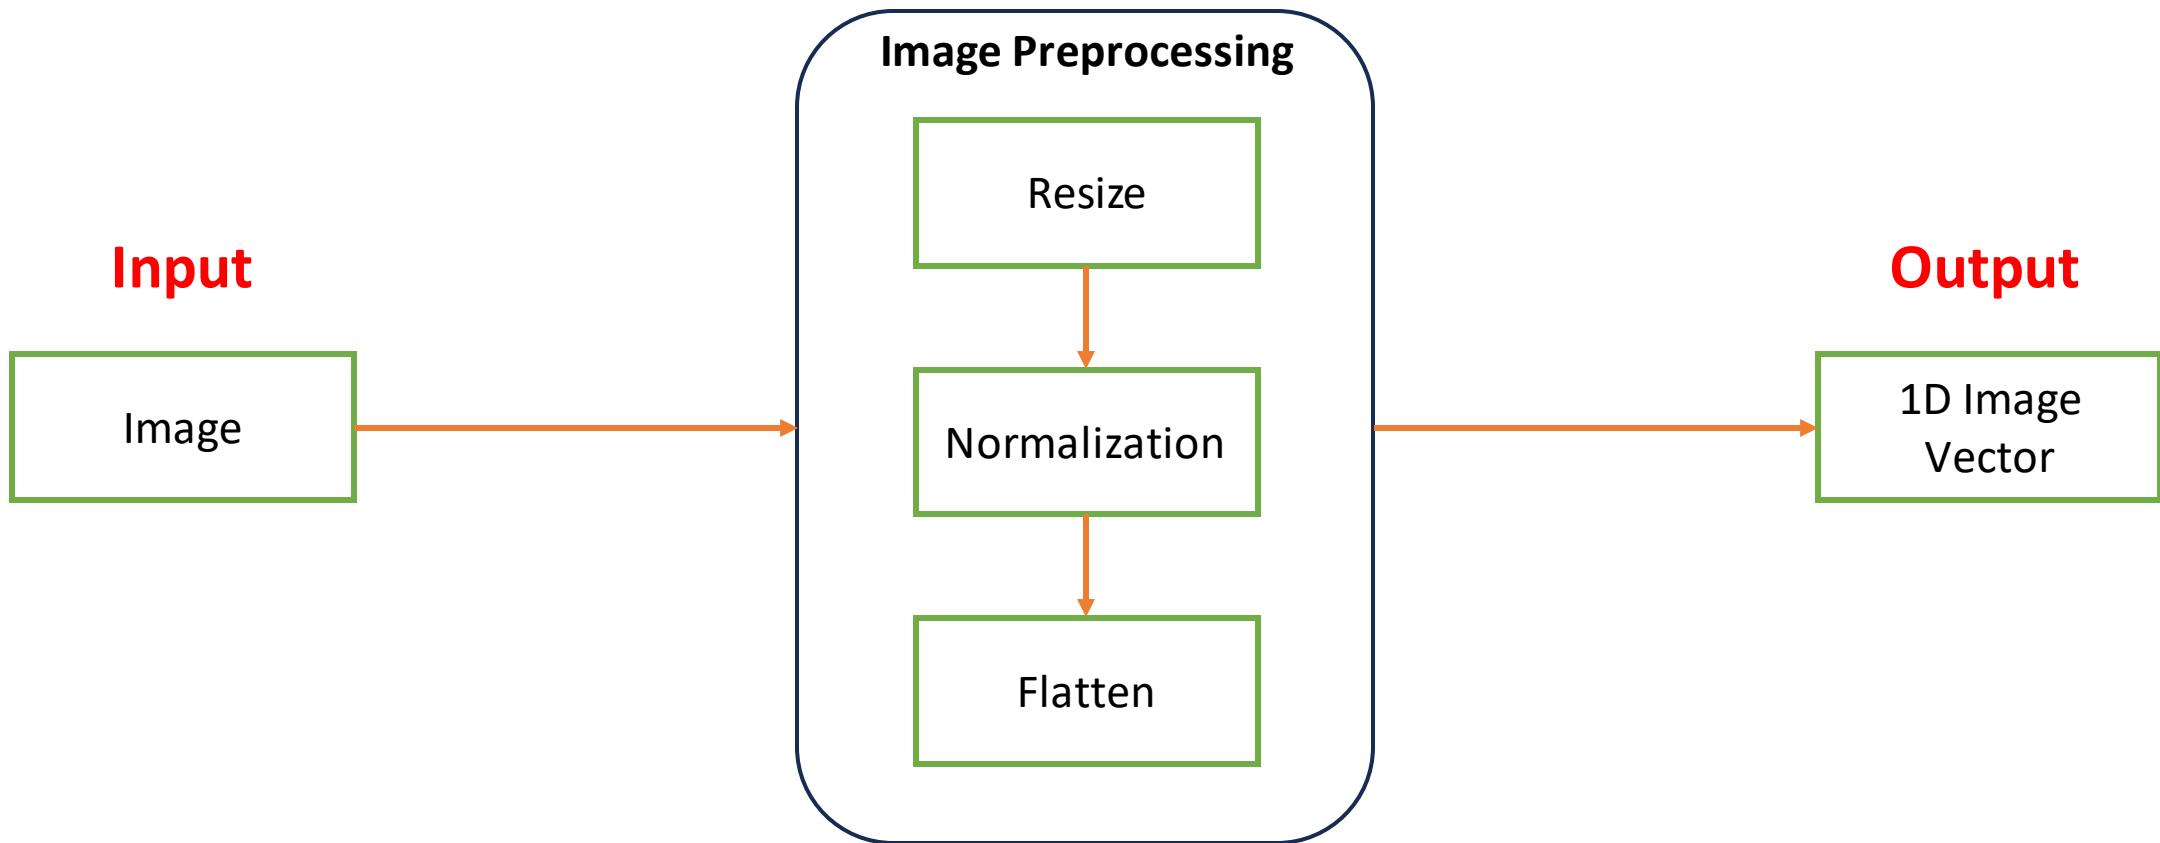
# Baseline

## ❖ Project Pipeline



# Baseline

## ❖ Step 3: Preprocess image



# Baseline

## ❖ Step 3.1: Resize image

```
cv2.resize(src, dsize)
```

Resize an image into desired (height, width).



```
1 def image_resize(images, target_size=(64, 64)):  
2     resized_image = cv2.resize(  
3         images,  
4         target_size  
5     )  
6  
7     return resized_image
```



Resize to 64x64



# Baseline

## ❖ Step 3.2: Normalize image

**Image Normalization:** Transforms the pixel values of image into a common scale that might boost the performance of system.

**Z-score Normalization (Standardization):** A technique that transforms data so that it has a mean of 0 and a std of 1.

$$X_{normalized} = \frac{X - mean}{std}$$

[[[ 4 6 9]	[[[-1.20057116 -1.15850123 -0.97975352]
[ 0 6 0]	[-1.25822794 -1.15850123 -1.11004419]
[ 1 10 0]	[-1.24381375 -1.09741939 -1.11004419]
...	...
[ 0 8 0]	[-1.25822794 -1.12796031 -1.11004419]
[ 4 7 8]	[-1.20057116 -1.14323077 -0.99423026]
[ 3 1 21]]	[-1.21498535 -1.23485353 -0.80603264]]
[[ 0 2 2]	[[ -1.25822794 -1.21958307 -1.08109071]
[ 17 26 11]	[-1.01318662 -0.85309204 -0.95080004]
[ 26 40 11]	[-0.88345886 -0.6393056 -0.95080004]
...	...
[ 10 23 2]	[-1.11408599 -0.89890342 -1.08109071]
[ 3 7 5]	[-1.21498535 -1.14323077 -1.03766049]
[ 0 1 14]]	[-1.25822794 -1.23485353 -0.90736982]]
[[ 0 2 1]	[[ -1.25822794 -1.21958307 -1.09556745]
[ 38 47 26]	[-0.71048852 -0.53241238 -0.73364894]
[ 49 65 24]	[-0.55193237 -0.25754411 -0.76260242]
...	...
[ 60 80 36]	[-0.39337623 -0.02848721 -0.58888154]
[ 11 20 5]	[-1.09967179 -0.9447148 -1.03766049]
[ 0 3 6]]	[-1.25822794 -1.20431261 -1.02318375]]

...

...

# Baseline

## ❖ Step 3.2: Normalize image



```
1 def calculate_mean_std(images):
2     mean = np.mean(images, axis=(0, 1, 2))
3     std = np.std(images, axis=(0, 1, 2))
4
5     return mean, std
6
7 def image_std_normalize(images, mean, std):
8     normalized_image = (images - mean) / std
9
10    return normalized_image
```

**Z-score Normalization (Standardization):** A technique that transforms data so that it has a mean of 0 and a std of 1.

$$X_{normalized} = \frac{X - mean}{std}$$

# Baseline

## ❖ Step 3.3: Flatten image



```
1 def image_flatten(images, is_batch=False):
2     # For flatten a bunch of images
3     if is_batch:
4         flattened_image = images.reshape(
5             images.shape[0], -1
6         )
7     # For flatten a single image
8     else:
9         flattened_image = images.reshape(-1)
10
11 return flattened_image
```

`ndarray.reshape(shape, order='C')`

Returns an array containing the same data with a new shape.

Refer to [numpy.reshape](#) for full documentation.

See also

[numpy.reshape](#)

equivalent function

```
6 normalized_image = image_std_normalize(
7     image_resized,
8     mean, std
9 )
10 flattened_image = image_flatten(normalized_image)
11
12 print('Shape before flattening: ', normalized_image.shape)
13 print('Shape after flattening: ', flattened_image.shape)
```

Shape before flattening: (64, 64, 3)

Shape after flattening: (12288,)

# Baseline

## ❖ Step 3.4: Final image preprocessing function



```
1 def preprocess_batches(images):
2     resized_images = [
3         image_resize(image) for image in src_images
4     ]
5     images_arr = np.array(resized_images)
6     mean, std = calculate_mean_std(images_arr)
7     normalized_images = image_std_normalize(
8         images_arr,
9         mean, std
10    )
11    flattened_images = image_flatten(
12        normalized_images,
13        is_batch=True
14    )
15
16    return flattened_images, mean, std
```



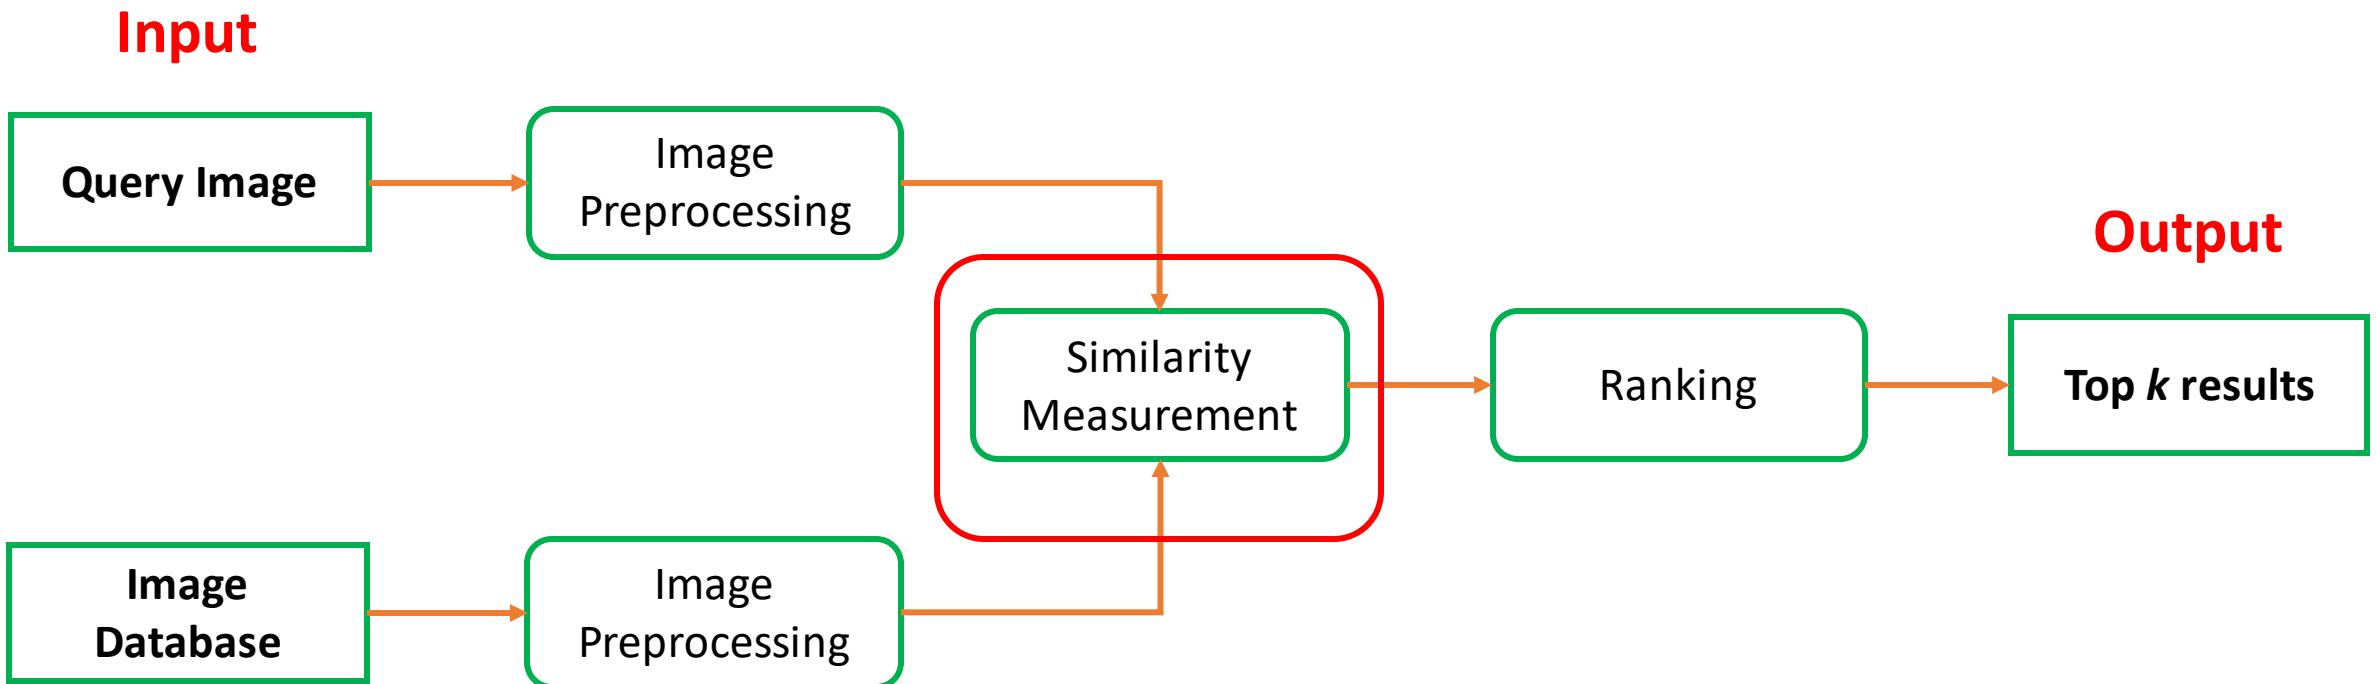
```
1 def preprocess_query(image):
2     resized_image = image_resize(image)
3     normalized_image = image_std_normalize(
4         resized_image,
5         mean, std
6     )
7     flattened_image = image_flatten(normalized_image)
8
9     return flattened_image
```

Function to process the whole dataset

Function to process a single image

# Baseline

## ❖ Project Pipeline

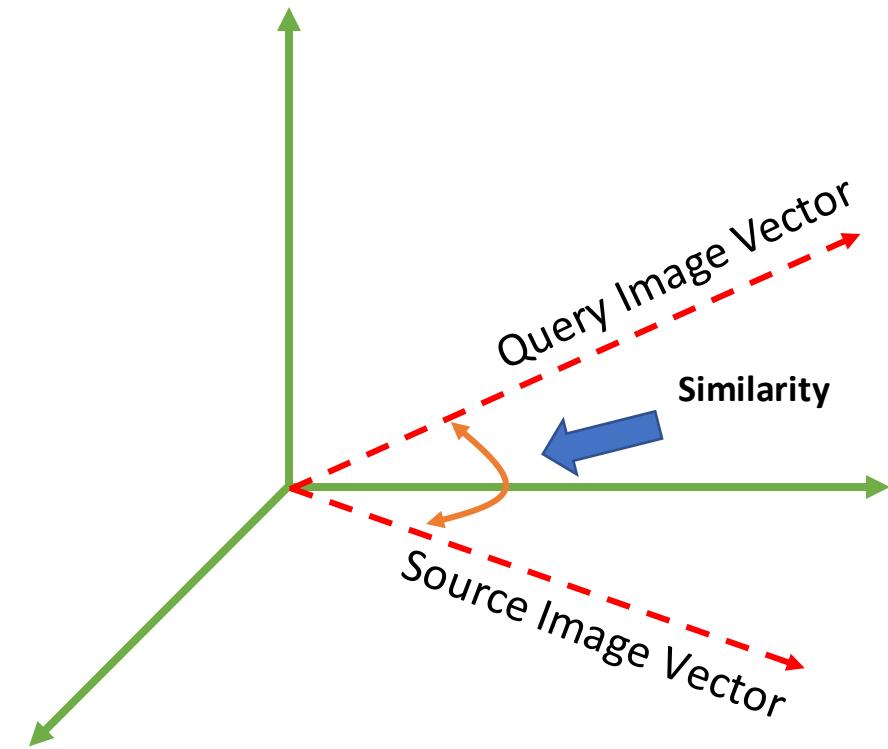


# Baseline

## ❖ Step 4: Define similarity metrics

**Similarity Calculation:** Define a function to find the similarity between two vectors. In this project, we will define 4 metrics:

1. Mean Absolute Error (L1 Loss)
2. Mean Squared Error (L2 Loss)
3. Cosine Similarity
4. Correlation Coefficient



# Baseline

## ❖ Step 4: Mean Absolute Error



```
1 def mean_absolute_error(query_vector, src_vectors):  
2     abs_diff = np.abs(src_vectors - query_vector)  
3  
4     mae = np.mean(abs_diff, axis=1)  
5  
6     return mae
```

Given two 1D image vectors that have  $n$  elements, the Mean Absolute Error is calculated as follows:

$$MAE = \frac{1}{n} \sum_{i=1}^n |a_i - b_i|$$

# Baseline

## ❖ Step 4: Mean Squared Error



```
1 def mean_squared_error(query_vector, src_vectors):  
2     squared_diff = (src_vectors - query_vector) ** 2  
3     mse = np.mean(squared_diff, axis=1)  
4     return mse
```

Given two 1D image vectors that have  $n$  elements, the Mean Squared Error is calculated as follows:

$$MSE = \frac{1}{n} \sum_{i=1}^n (a_i - b_i)^2$$

# Baseline

## ❖ Step 4: Cosine Similarity



```
1 def cosine_similarity(query_vector, src_vectors):
2     query_norm = np.linalg.norm(query_vector)
3     normalized_query = query_vector / query_norm
4     src_norms = np.linalg.norm(src_vectors, axis=1)
5     normalized_src = src_vectors / src_norms[:, np.newaxis]
6
7     cosine_similarity = np.dot(
8         normalized_src,
9         normalized_query
10    )
11
12     return cosine_similarity
```

Given two 1D image vectors that have  $n$  elements, the Cosine Similarity is calculated as follows:

$$\text{cosine\_similarity} = \frac{a \cdot b}{\|a\| \|b\|}$$

# Baseline

## ❖ Step 4: Correlation Coefficient



```
1 def correlation_coefficient(query_vector, src_vectors):
2     corr_coef = np.corrcoef(
3         query_vector,
4         src_vectors
5     )[:-1, -1]
6
7     return corr_coef
```

`numpy.corrcoef(x, y=None, rowvar=True, bias=<no value>, ddof=<no value>, *, dtype=None)` [\[source\]](#)

Return Pearson product-moment correlation coefficients.

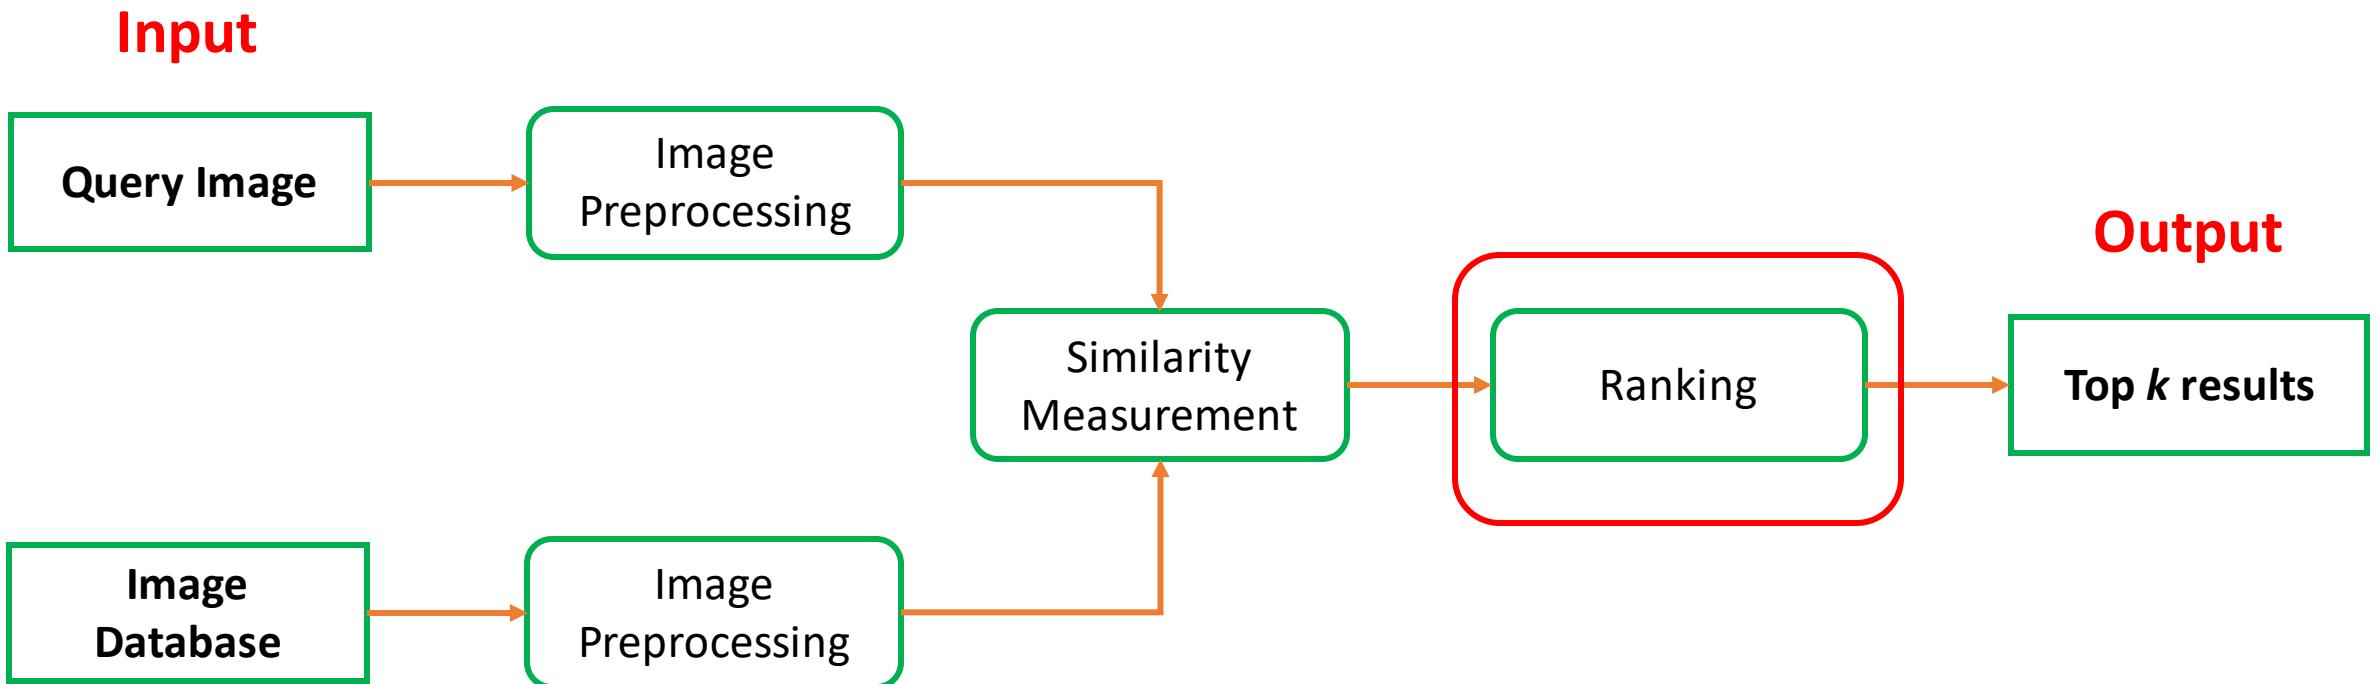
Please refer to the documentation for `cov` for more detail. The relationship between the correlation coefficient matrix,  $R$ , and the covariance matrix,  $C$ , is

$$R_{ij} = \frac{C_{ij}}{\sqrt{C_{ii}C_{jj}}}$$

The values of  $R$  are between -1 and 1, inclusive.

# Baseline

## ❖ Project Pipeline



# Baseline

## ❖ Step 5: Ranking

```
1 def ranking(  
2     preprocessed_query_image,  
3     preprocessed_src_images,  
4     top_k=10  
5 ):  
6     scores = cosine_similarity(  
7         preprocessed_query_image,  
8         preprocessed_src_images  
9     )  
10    ranked_list = np.argsort(scores)[::-1][:top_k]  
11    scores = scores[ranked_list]  
12  
13    return ranked_list, scores
```

# Baseline

## ❖ Ranking program

●

```
1 query_image_paths = [  
2     '/content/images_mr/1.jpg'  
3 ]  
4 top_k = 10  
5  
6 for query_image_path in query_image_paths:  
7     query_image = cv2.imread(query_image_path, 1)  
8     query_image = cv2.cvtColor(query_image, cv2.COLOR_BGR2RGB)  
9     preprocessed_query_image = preprocess_query(query_image)  
10  
11 ranked_list, scores = ranking(  
12     preprocessed_query_image,  
13     preprocessed_src_images,  
14     top_k  
15 )  
16  
17 print('Query Image')  
18 plt.figure(figsize=(3, 3))  
19 plt.imshow(query_image)  
20 plt.axis('off')  
21 plt.show()  
22 print(f'Top {top_k} results')  
23 for idx in range(len(ranked_list)):  
24     src_image_idx = ranked_list[idx]  
25     similarity_score = scores[idx]  
26     plt.figure(figsize=(3, 3))  
27     plt.imshow(src_images[src_image_idx])  
28     plt.title(f'Similarity: {similarity_score}', fontsize=10)  
29     plt.axis('off')  
30     plt.show()
```

Similarity: 1.0



Similarity: 0.6072830969370191



Similarity: 0.5779375079555944

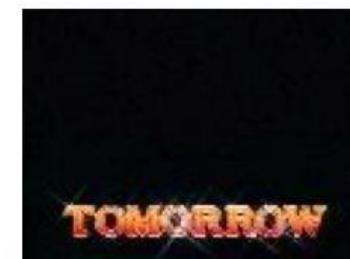
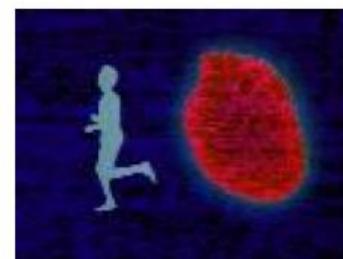


# Baseline

## ❖ Ranking program: How to visualize



**Matplotlib:** A library for creating static, animated, and interactive visualizations in Python.



# Baseline

## ❖ Ranking program: How to visualize

### matplotlib.pyplot.figure

```
matplotlib.pyplot.figure(num=None, figsize=None, dpi=None, *,  
facecolor=None, edgecolor=None, frameon=True, FigureClass=<class  
'matplotlib.figure.Figure'>, clear=False, **kwargs)
```

Create a new figure, or activate an existing figure.

[source]



```
1 plt.figure()  
2 plt.show()  
3  
4 # <Figure size 640x480 with 0 Axes>
```

### matplotlib.pyplot.show

```
matplotlib.pyplot.show(*, block=None)
```

Display all open figures.

[source]



```
1 plt.figure(figsize=(12, 12))  
2 plt.show()  
3  
4 # <Figure size 1200x1200 with 0 Axes>
```

# Baseline

## ❖ Ranking program: How to visualize



```
1 plt.figure(figsize=(4, 4))
2 image_path = 'images_mr/1.jpg'
3 image = cv2.imread(image_path)
4 image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
5 plt.imshow(image)
6 plt.show()
```



## matplotlib.pyplot.imshow

```
matplotlib.pyplot.imshow(X, cmap=None, norm=None, *, aspect=None,
interpolation=None, alpha=None, vmin=None, vmax=None, origin=None,
extent=None, interpolation_stage=None, filternorm=True, filterrad=4.0,
resample=None, url=None, data=None, **kwargs)
```

[source]

Display data as an image, i.e., on a 2D regular raster.

# Baseline

## ❖ Ranking program: How to visualize



```
1 plt.figure(figsize=(4, 4))  
2 image_path = 'images_mr/1.jpg'  
3 image = cv2.imread(image_path)  
4 image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)  
5 plt.imshow(image)  
6 plt.axis('off')  
7 plt.title(f'Image of {image_path}')  
8 plt.show()
```

Image of images\_mr/1.jpg



# Baseline

## ❖ Baseline results



Ranking



top\_k=4 results

Similarity: 0.9999999999999993



Similarity: 0.625198062681602



Similarity: 0.5972609035600188

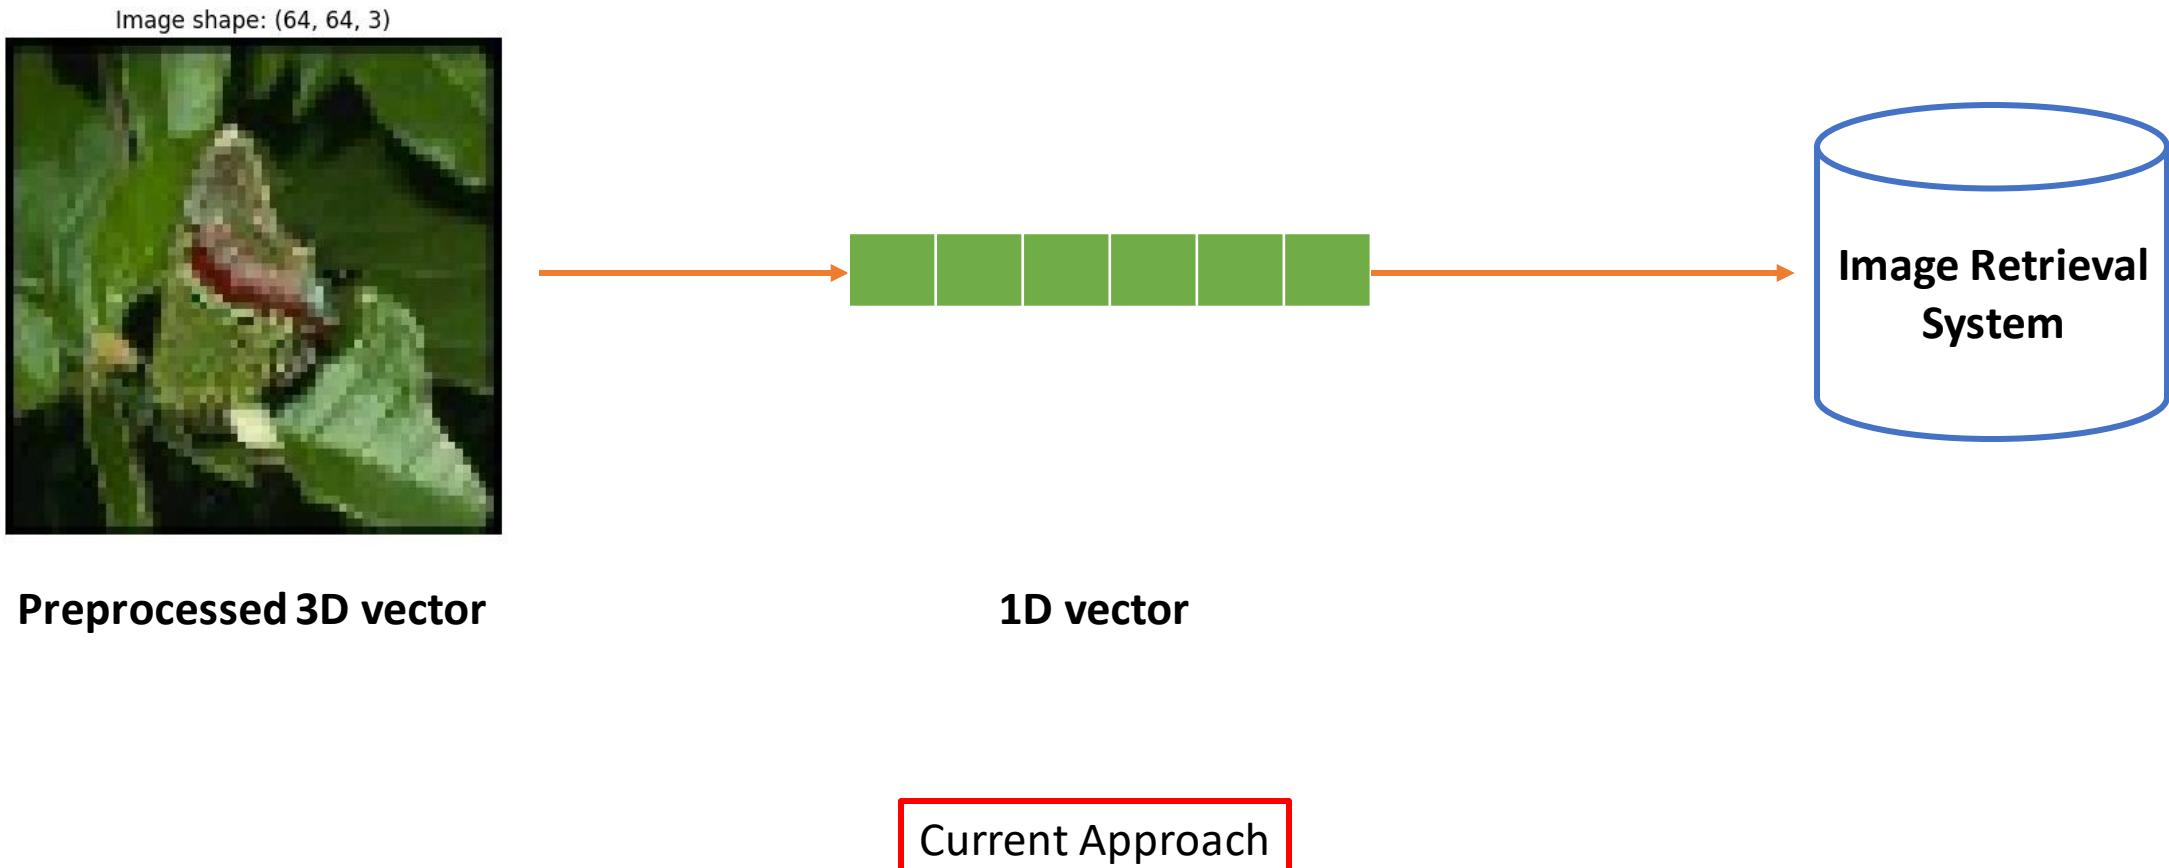


Similarity: 0.5932664796583871



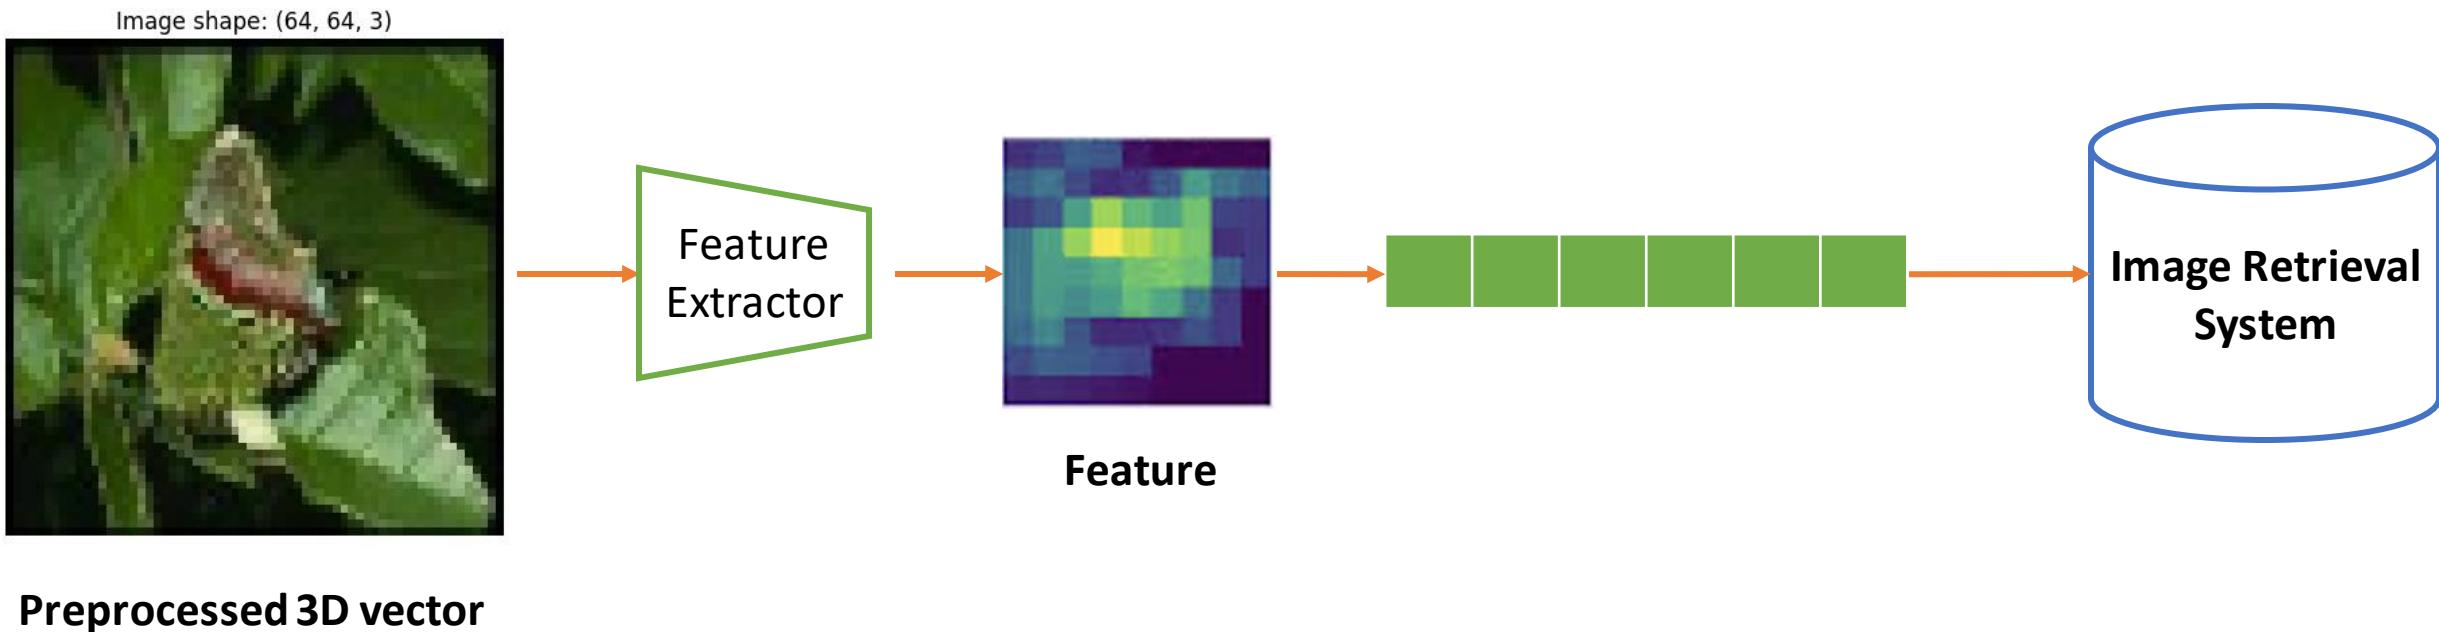
# Apply pre-trained ViT

## ❖ Getting Started



# Apply pre-trained ViT

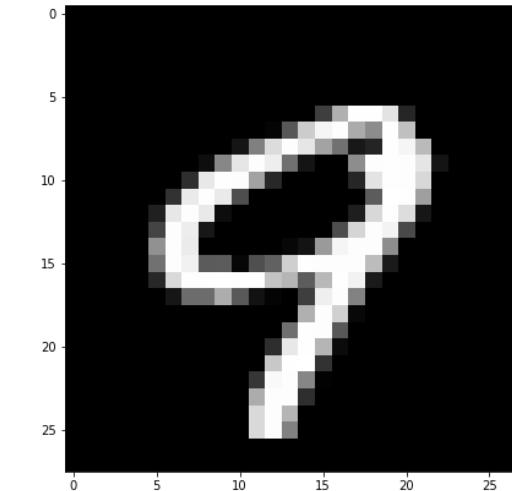
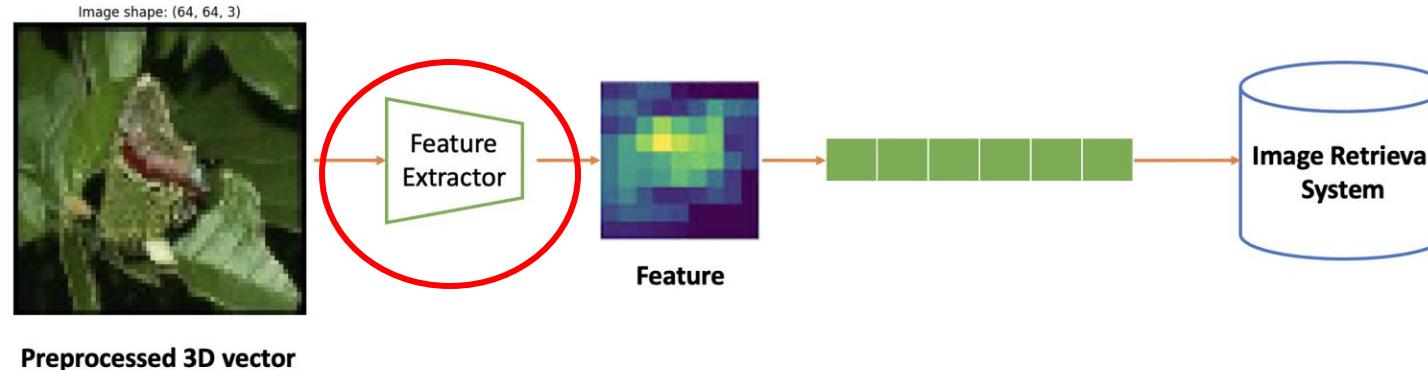
## ❖ Getting Started



Modern approach with the use of ML/DL model

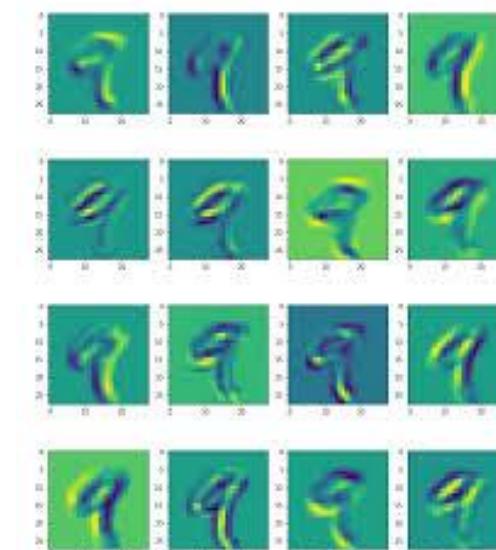
# Apply pre-trained ViT

## ❖ Getting Started



**Feature Extractor:** A component (typically a network) that is automatically identifying and extracting meaningful features from raw image data. Some well-known extractor are:

- **Digital Image Processing:** HOG, SIFT...
- **Traditional Machine Learning:** Bag of Visual Words, PCA...
- **Deep Learning:** CNNs, Autoencoder, ViT...



Raw Image

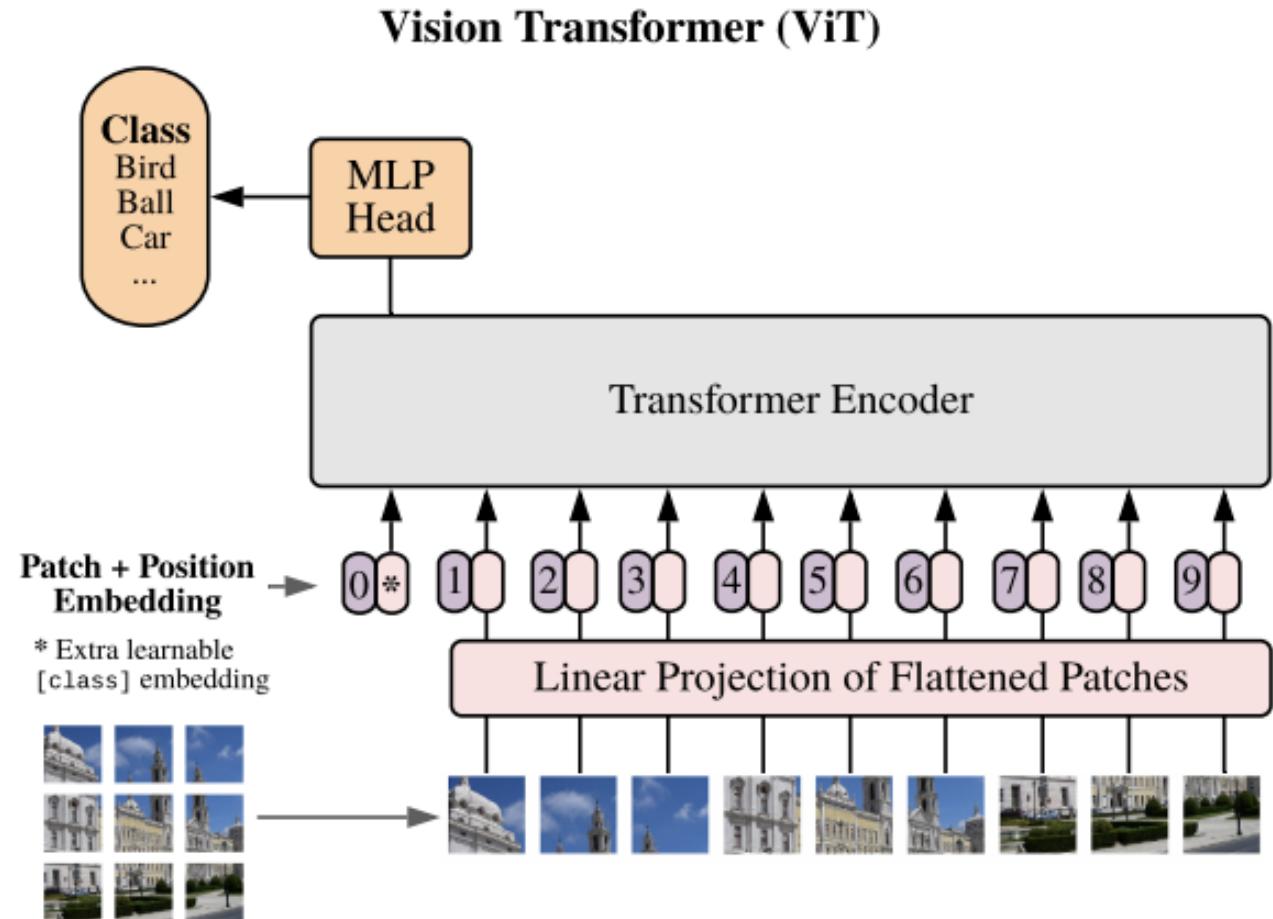
Feature Maps  
(resulted from  
Feature  
Extractor)

# Apply pre-trained ViT

## ❖ Introduction

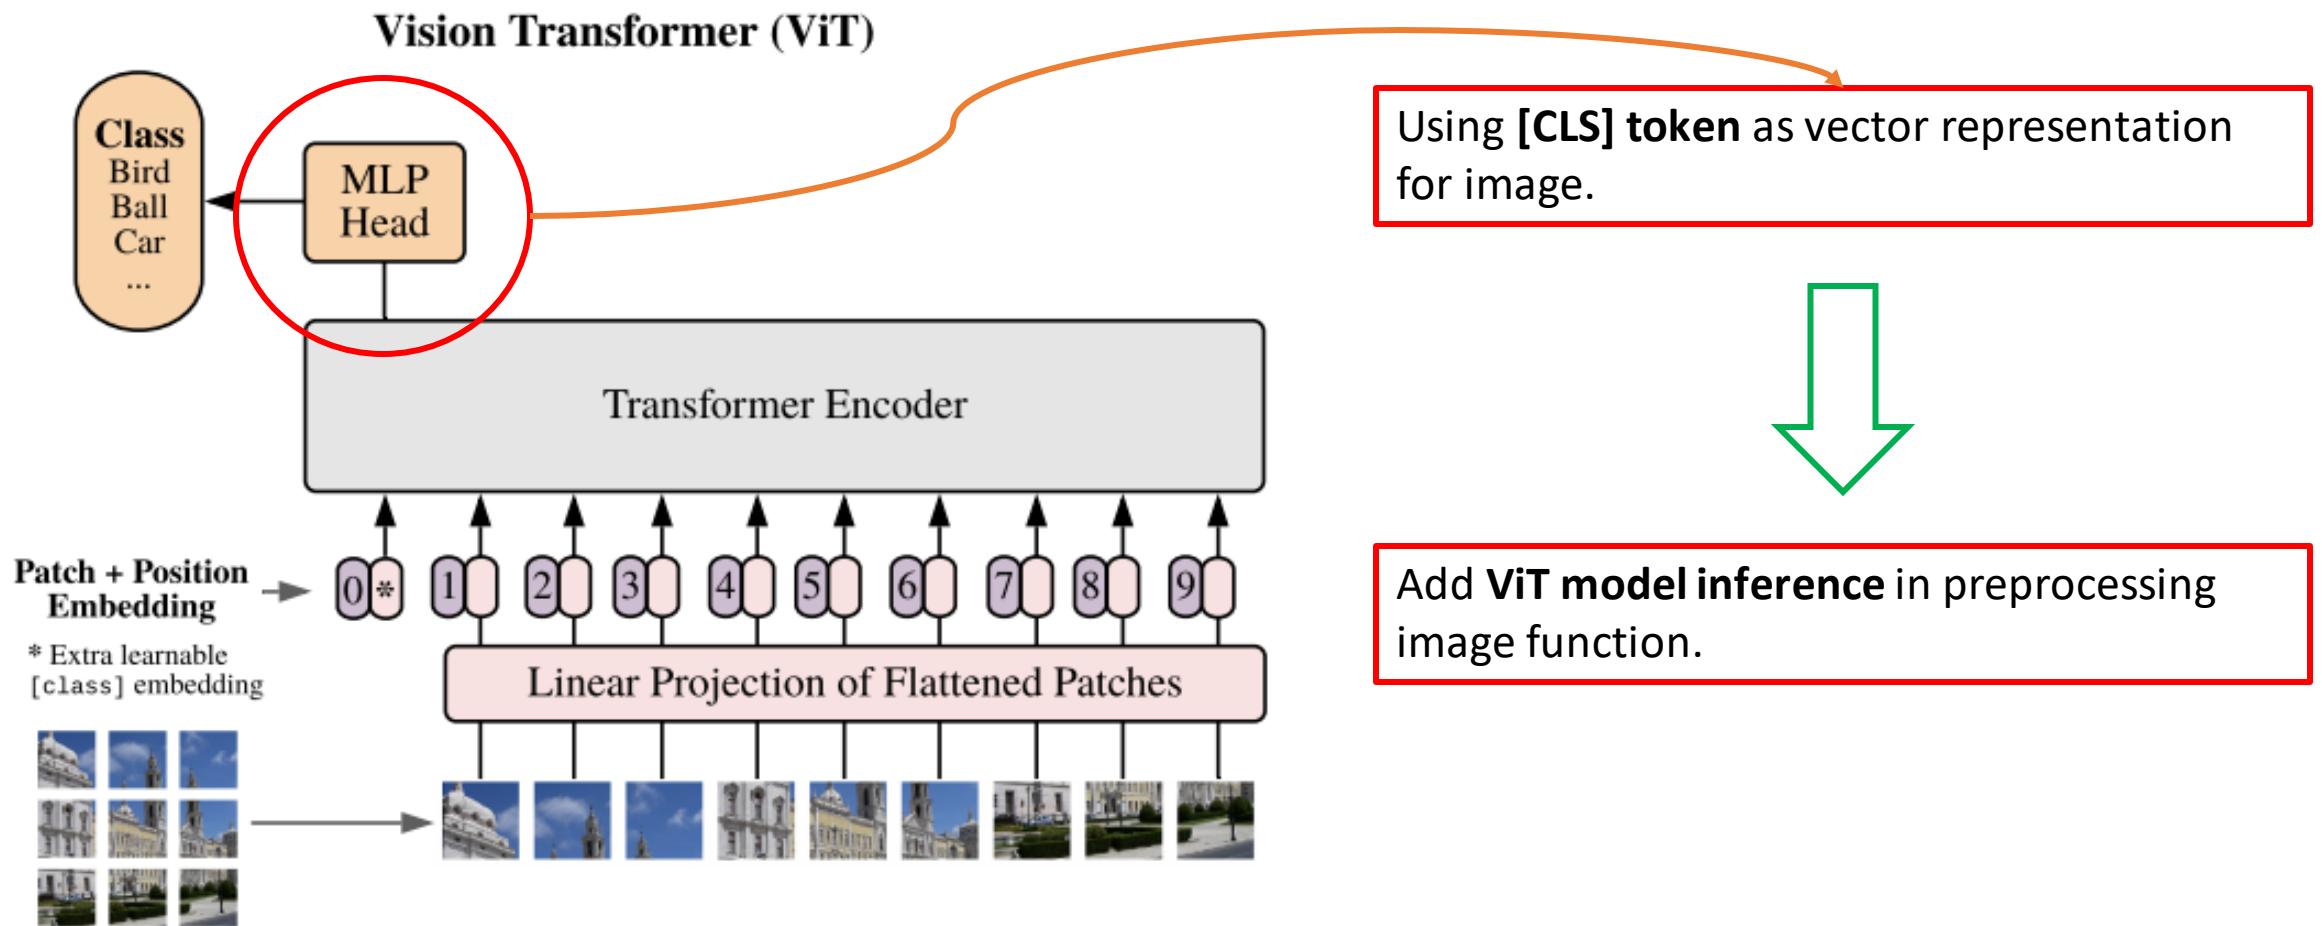
**Vision Transformer (ViT):** A deep learning architecture designed for image classification tasks. ViT deviates from traditional CNNs by using a transformer-based architecture.

❖ Paper: <https://arxiv.org/pdf/2010.11929.pdf>



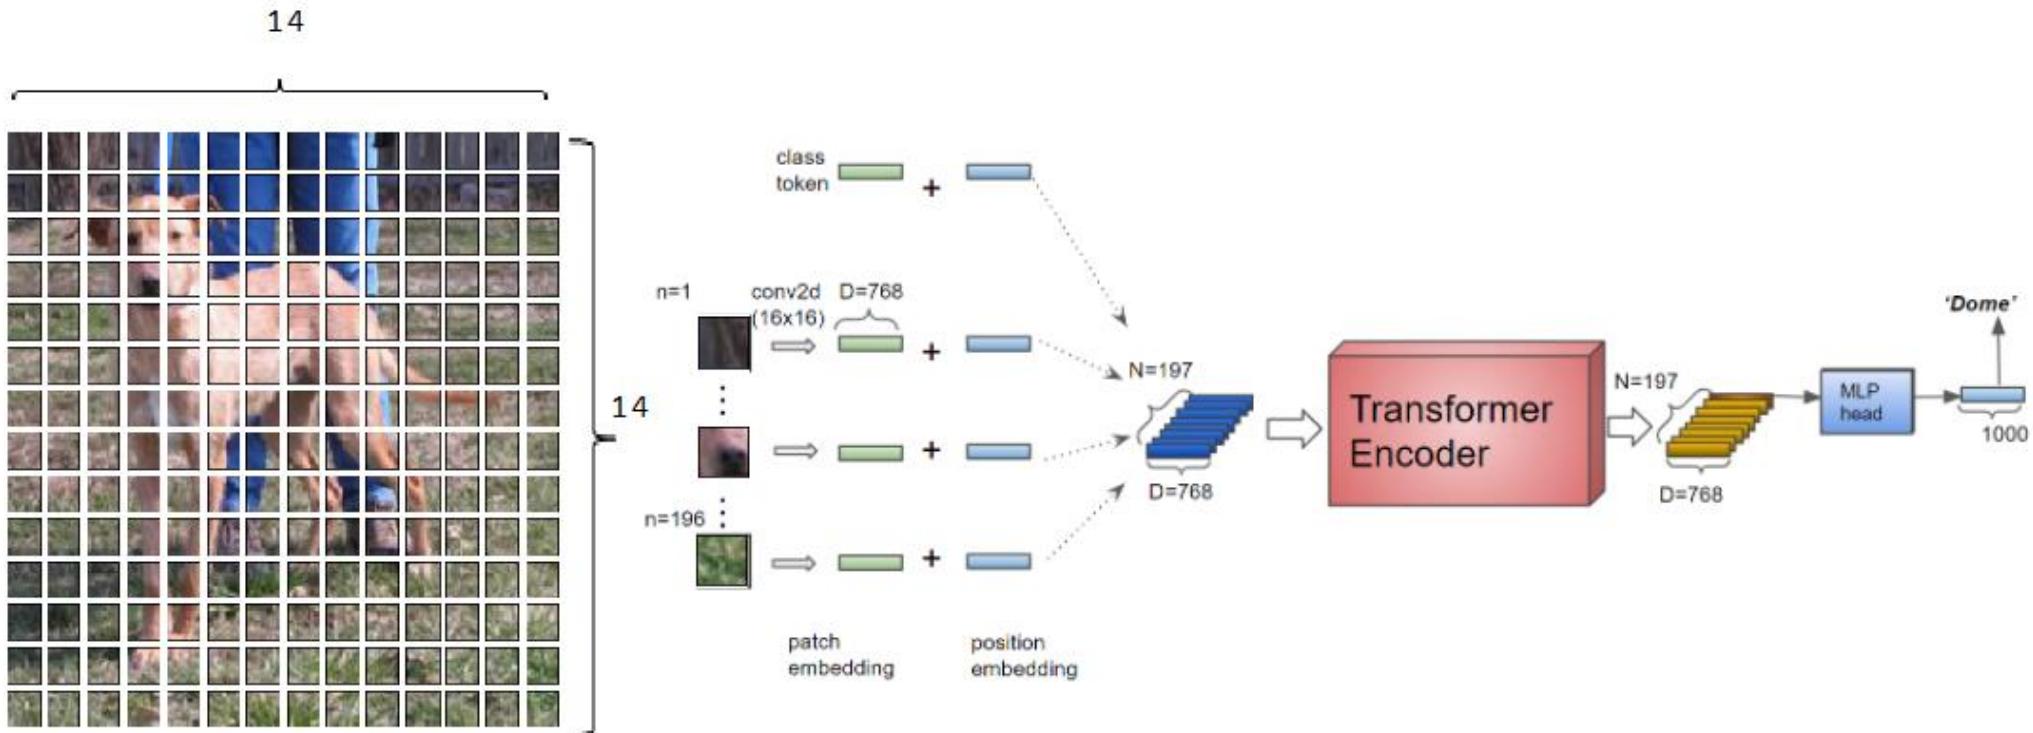
# Apply pre-trained ViT

## ❖ Idea



# Apply pre-trained ViT

## ❖ Idea



Using **[CLS]** token as vector representation for image.

# Apply pre-trained ViT

## ❖ Step 1: Import libraries and get filenames

### 1. Import necessary libraries

```
● ○ ●  
1 import os  
2 import numpy as np  
3 import cv2  
4 import matplotlib.pyplot as plt  
5 import torch  
6 from transformers import ViTImageProcessor,  
    ViTForImageClassification
```

### 2. Get image filenames

```
● ○ ●  
1 # Declare source images directory path  
2 dataset_dir = 'images_mr'  
3  
4 # Get all filenames inside "dataset_dir"  
5 image_filenames = os.listdir(dataset_dir)
```



```
1 !pip install transformers==4.31.0 -q
```

```
1 image_filenames
```

```
'9385.jpg',  
'5052.jpg',  
'6271.jpg',  
'5362.jpg',  
'5852.jpg',  
'6164.jpg',  
'8962.jpg',  
'3608.jpg',  
'6438.jpg',  
'6670.jpg',  
'5018.jpg',  
'8697.jpg',
```

```
1 print(  
2     'Number of files: ',  
3     len(image_filenames)  
4 )
```

```
Number of files: 9908
```

**Note:** Turn on GPU before start notebook.

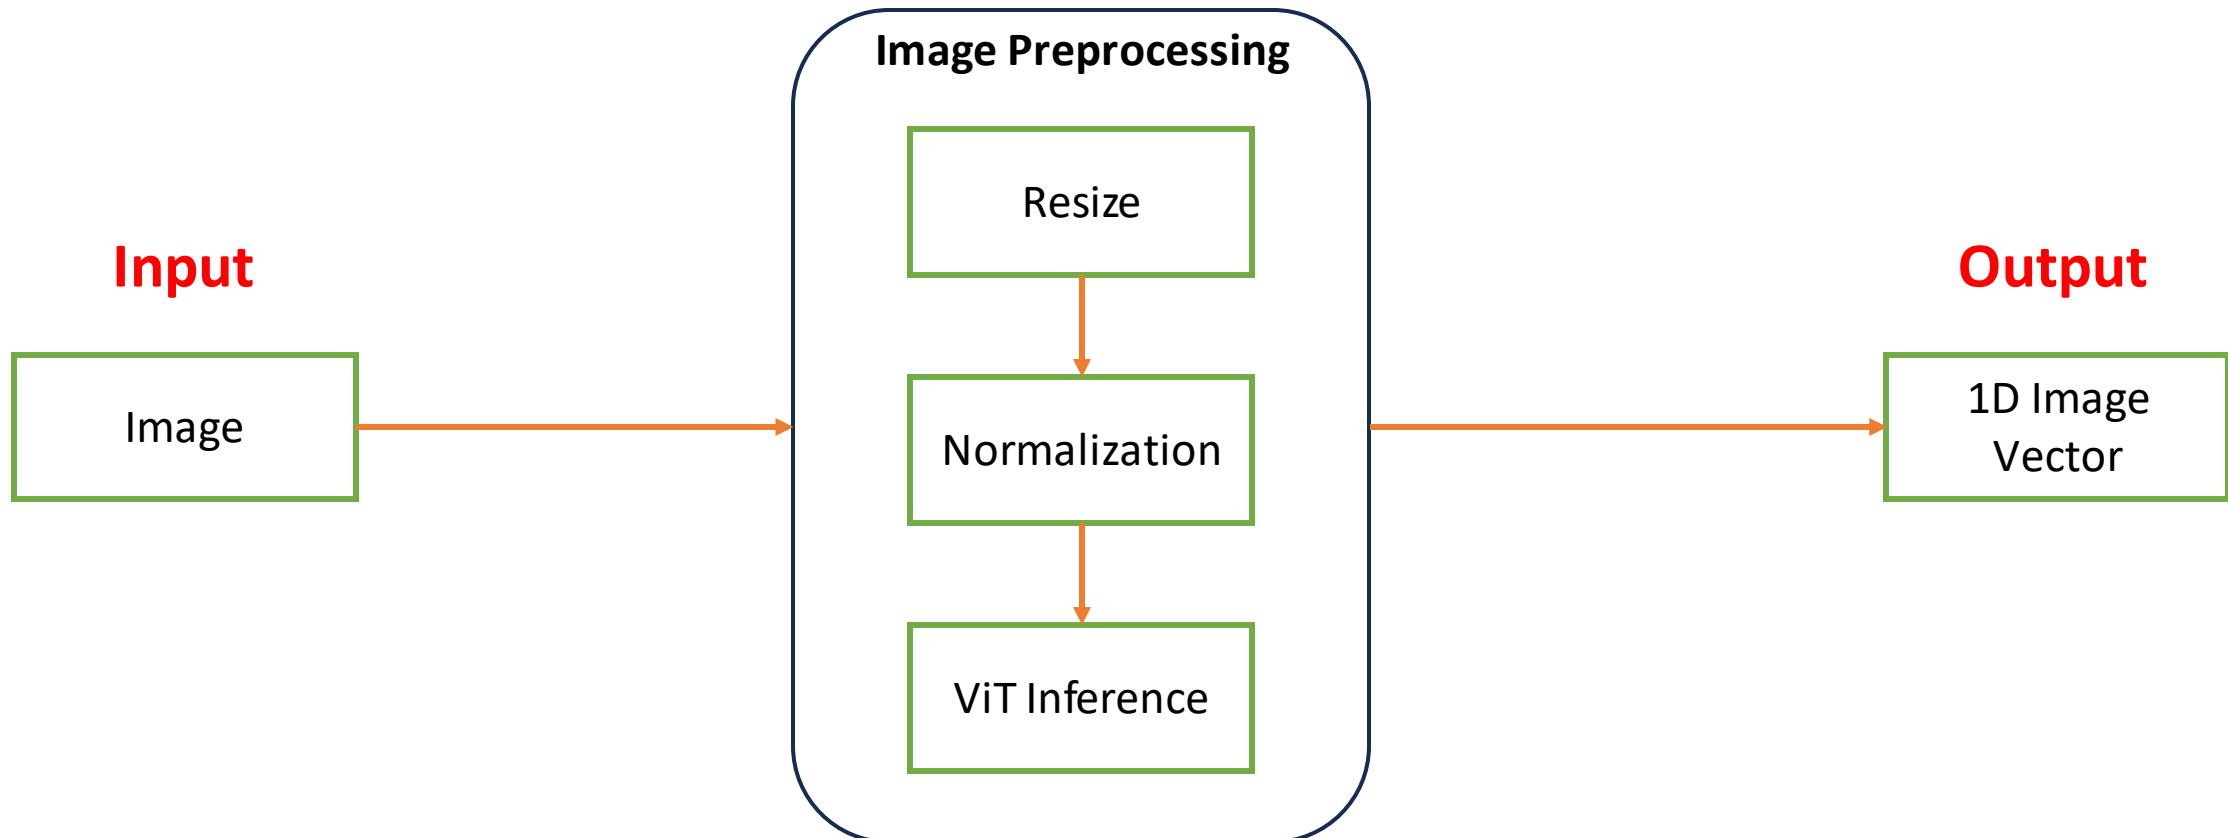
# Apply pre-trained ViT

## ❖ Step 2: Load image dataset

```
● ● ●  
1 dataset_dir = 'images_mr'  
2 image_filenames = os.listdir(dataset_dir)[:500]  
3 src_images = []  
4 for filename in image_filenames:  
5     filepath = os.path.join(  
6         dataset_dir,  
7         filename  
8     )  
9  
10    image = cv2.imread(filepath)  
11    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)  
12    src_images.append(image)
```

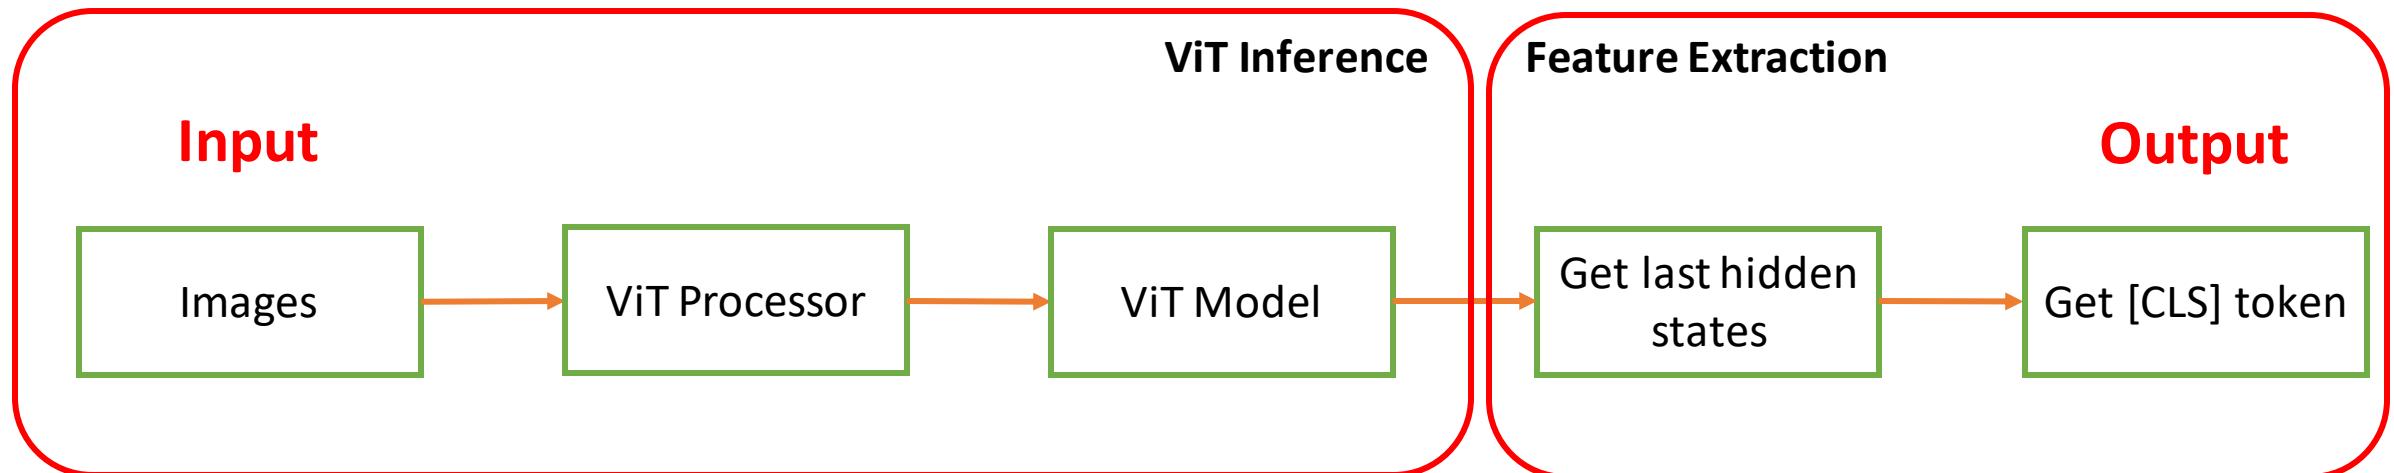
# Apply pre-trained ViT

## ❖ Step 3: Define preprocessing function



# Apply pre-trained ViT

## ❖ Step 3: Define preprocessing function



New preprocessing pipeline with the use of ViT

# Apply pre-trained ViT

## ❖ Step 3: Define preprocessing function



Pre-trained ViT can be taken from  
**Huggingface** library.

```
1 device = 'cuda' if torch.cuda.is_available() else 'cpu'  
2 processor = ViTImageProcessor.from_pretrained(  
3     "google/vit-base-patch16-224")  
4 model = ViTForImageClassification.from_pretrained(  
5     "google/vit-base-patch16-224").to(device)
```

# Apply pre-trained ViT

## ❖ Step 3: Define preprocessing function

```
● ● ●  
1 def preprocessing(images):  
2     inputs = processor(  
3         images,  
4         return_tensors='pt'  
5     ).to(device)  
6  
7     with torch.no_grad():  
8         output = model(  
9             **inputs,  
10            output_hidden_states=True  
11        ).hidden_states[-1][:, 0, :].detach().cpu().numpy()  
12  
13    return output
```

# Apply pre-trained ViT

## ❖ Step 4, 5: Define similarity and rank function



```
1 def mean_absolute_error(query_vector, src_vectors):
2     abs_diff = np.abs(src_vectors - query_vector)
3     mae = np.mean(abs_diff, axis=1)
4
5     return mae
6
7 def mean_squared_error(query_vector, src_vectors):
8     squared_diff = (src_vectors - query_vector) ** 2
9     mse = np.mean(squared_diff, axis=1)
10
11    return mse
12
13 def cosine_similarity(query_vector, src_vectors):
14     query_norm = np.linalg.norm(query_vector)
15     normalized_query = query_vector / query_norm
16     src_norms = np.linalg.norm(src_vectors, axis=1)
17     normalized_src = src_vectors / src_norms[:, np.newaxis]
18
19     cosine_similarity = np.dot(normalized_src, normalized_query)
20
21     return cosine_similarity
22
23 def correlation_coefficient(query_vector, src_vectors):
24     return np.corrcoef(query_vector, src_vectors)[-1, -1]
```



```
1 def ranking(
2     preprocessed_query_image,
3     preprocessed_src_images,
4     top_k=10
5 ):
6     scores = cosine_similarity(
7         preprocessed_query_image,
8         preprocessed_src_images
9     )
10    ranked_list = np.argsort(scores)[-1:top_k]
11    scores = scores[ranked_list]
12
13    return ranked_list, scores
```

# Apply pre-trained ViT

## ❖ ViT results

Query Image



Ranking  
→

top\_k=4 results

Similarity: 0.5659981966018677



Similarity: 0.5508416295051575



Similarity: 0.5374281406402588



Similarity: 0.5243573188781738



# Question

